

Building Audi's enterprise big data platform

DataWorks Summit 2018 - Berlin

Matthias Graunitz (AUDI AG, Germany)

Carsten Herbe (Audi Business Innovation GmbH, Germany)



WHO ARE WE?

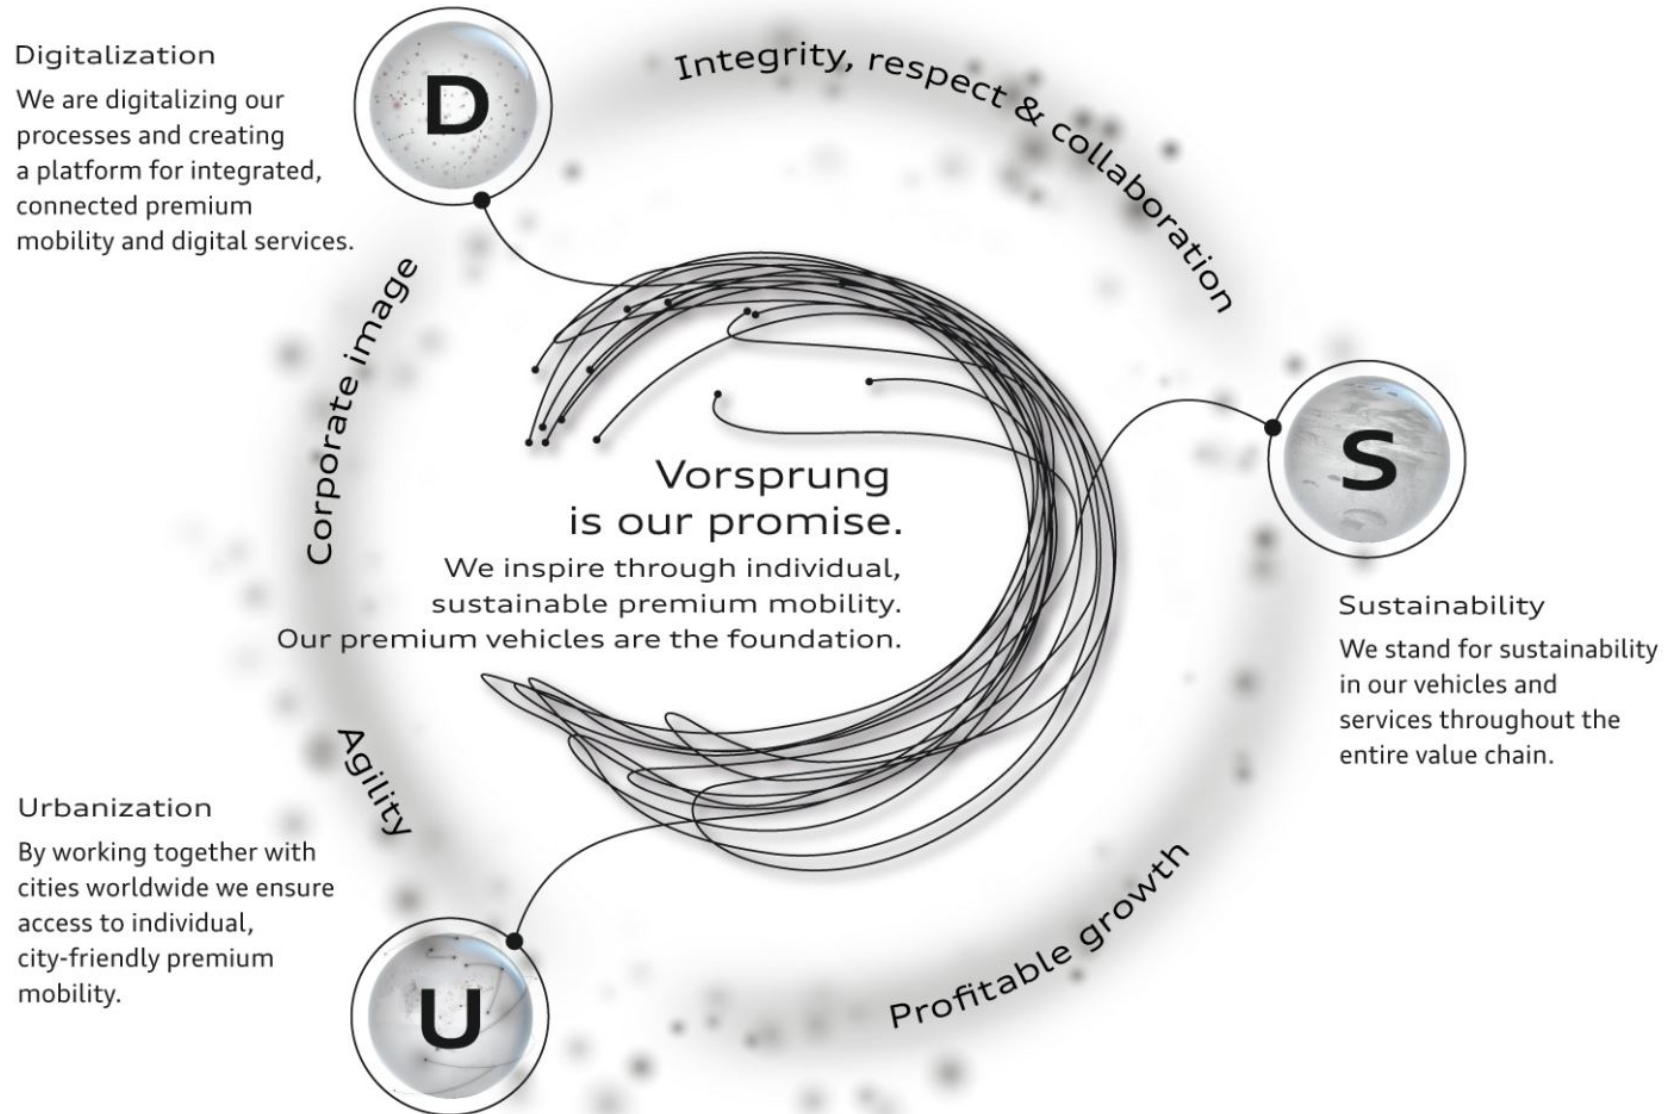
Audi Group

Audi, Lamborghini, Ducati and Italdesign



Vorsprung is our promise

Strategy 2025



Audi Business Innovation GmbH

...is the development, establishment, sales and operation of innovative concepts, products and services, as well the holding of shares in the field of future mobility.

Audi mobility
innovations



Audi on demand

Audi balanced
technologies



Audi e-gas

Audi customer
IT solutions



About us



Matthias Graunitz
AUDI AG

- » Center of Competence **Big Data & BI**
- » **Big Data** Architect
- » **10+ years** Data Warehousing & BI




Carsten Herbe
Audi Business Innovation GmbH

- » Data **Platform & Solution** Architecture
- » **Hadoop** since **2013**
- » **10+ years** Data Warehousing & BI


2 YEARS AGO...

STARTING BIG DATA AT AUDI

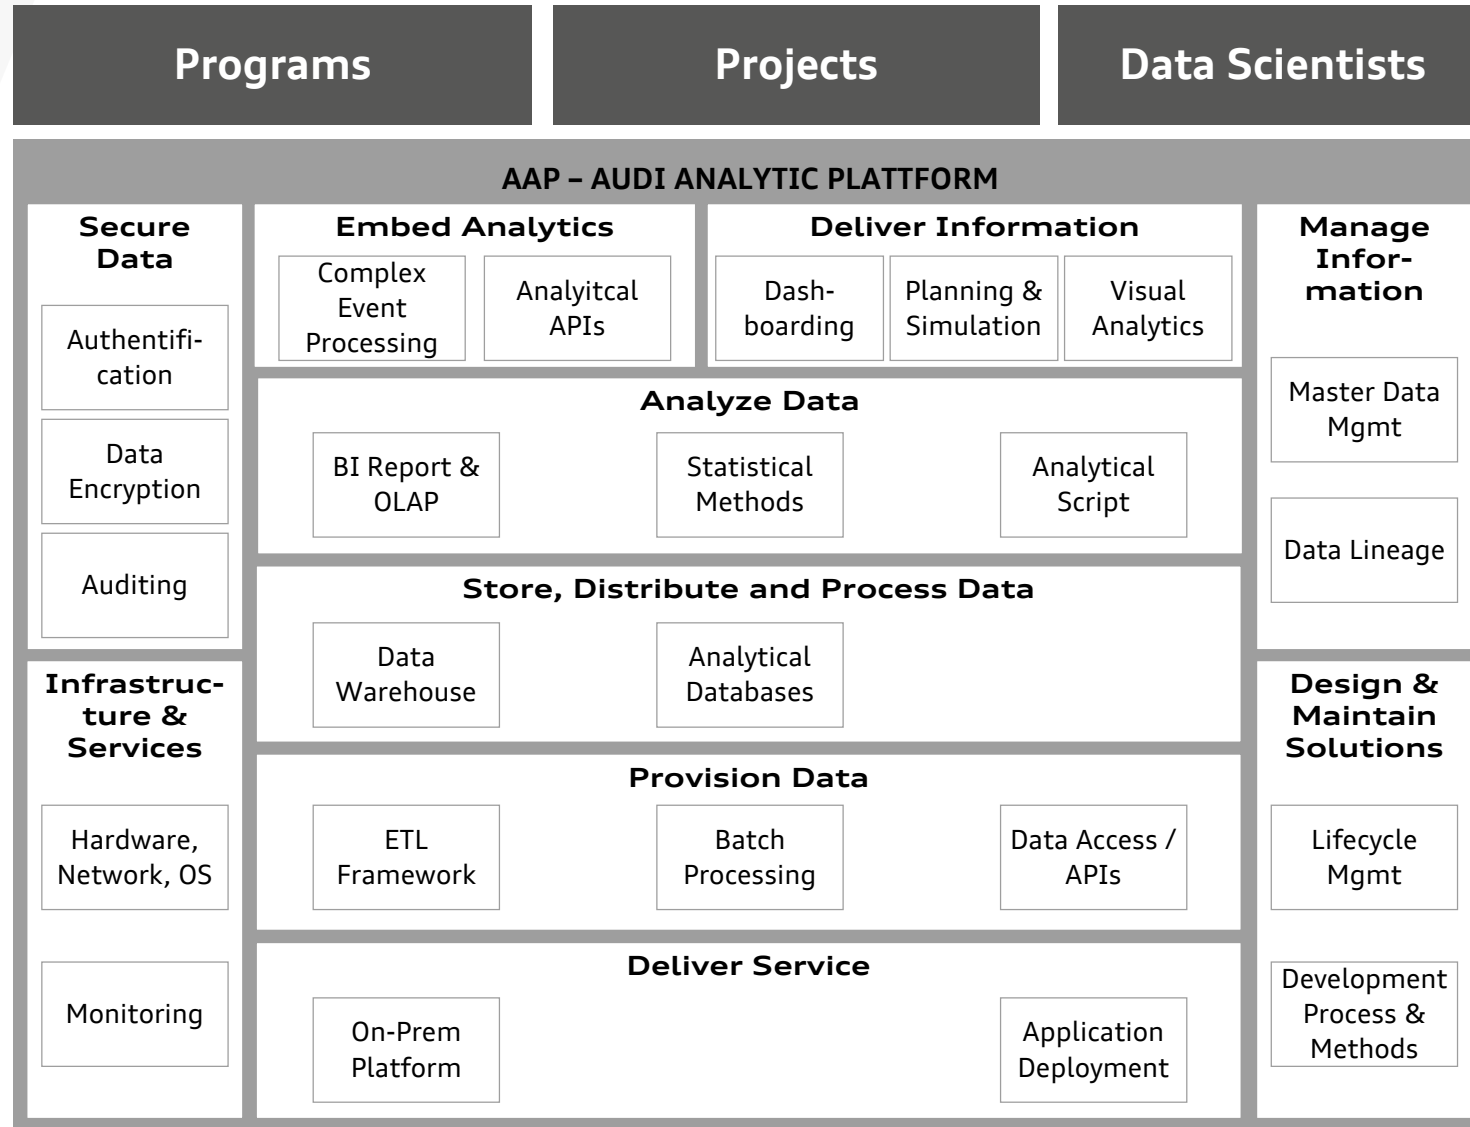
Analytical Capabilities by 2015




! Data Domains



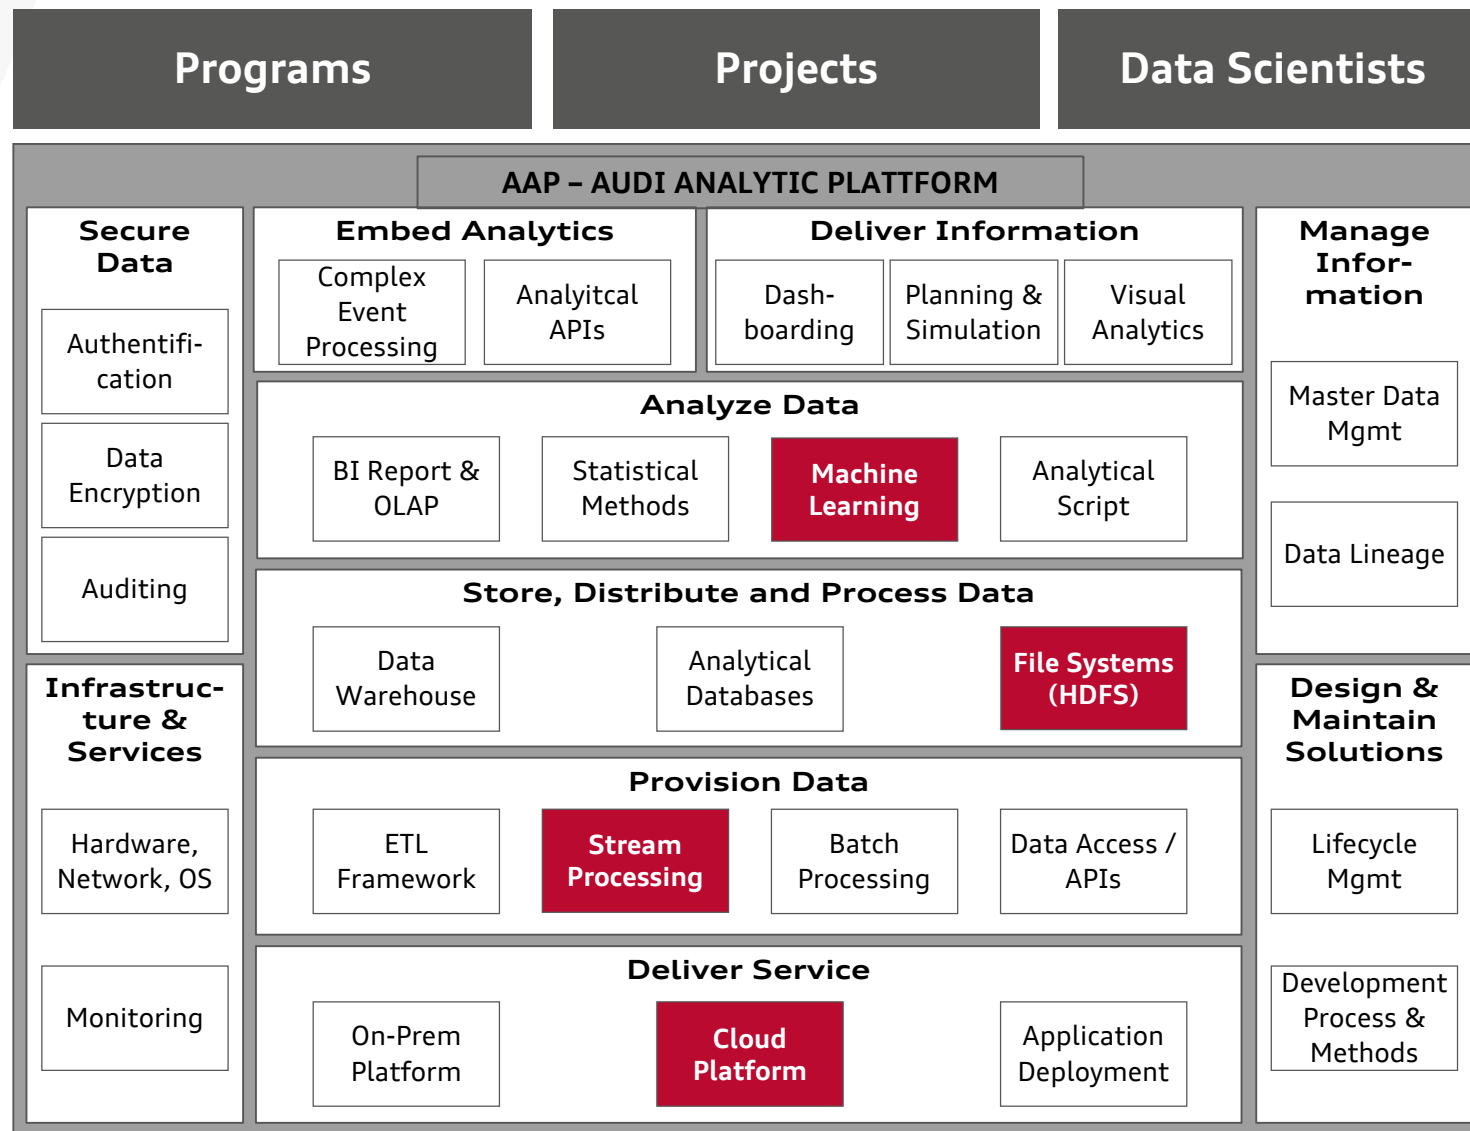
Analytical Capabilities by 2015




! Data Domains



Analytical Capabilities by 2015



! Data Domains



Our first Hadoop Cluster 2015



DEV

Hadoop	per node	Sum
# data nodes	1	4
RAM	128 GB	0,5 TB
Cores	24	96
HDD*	40 TB	160 TB


* Raw Capacity without replication and FS overhead!

Our first attempt to walk with Big Data Technologies





SCREWDRIVER ANALYSIS

Typische I.O.-Kurve mit 5 Stufen



- Diskrepanz zwischen W+ und W_O in Stufen 3&4 entdeckt
 - Optimierungspotential über 7 Minuten in 91 NIO-Fällen)
- Bug Schraubhersteller (fehlerhafte Übertragung Lösestufen) entdeckt

Skizze / Foto Schraubfall



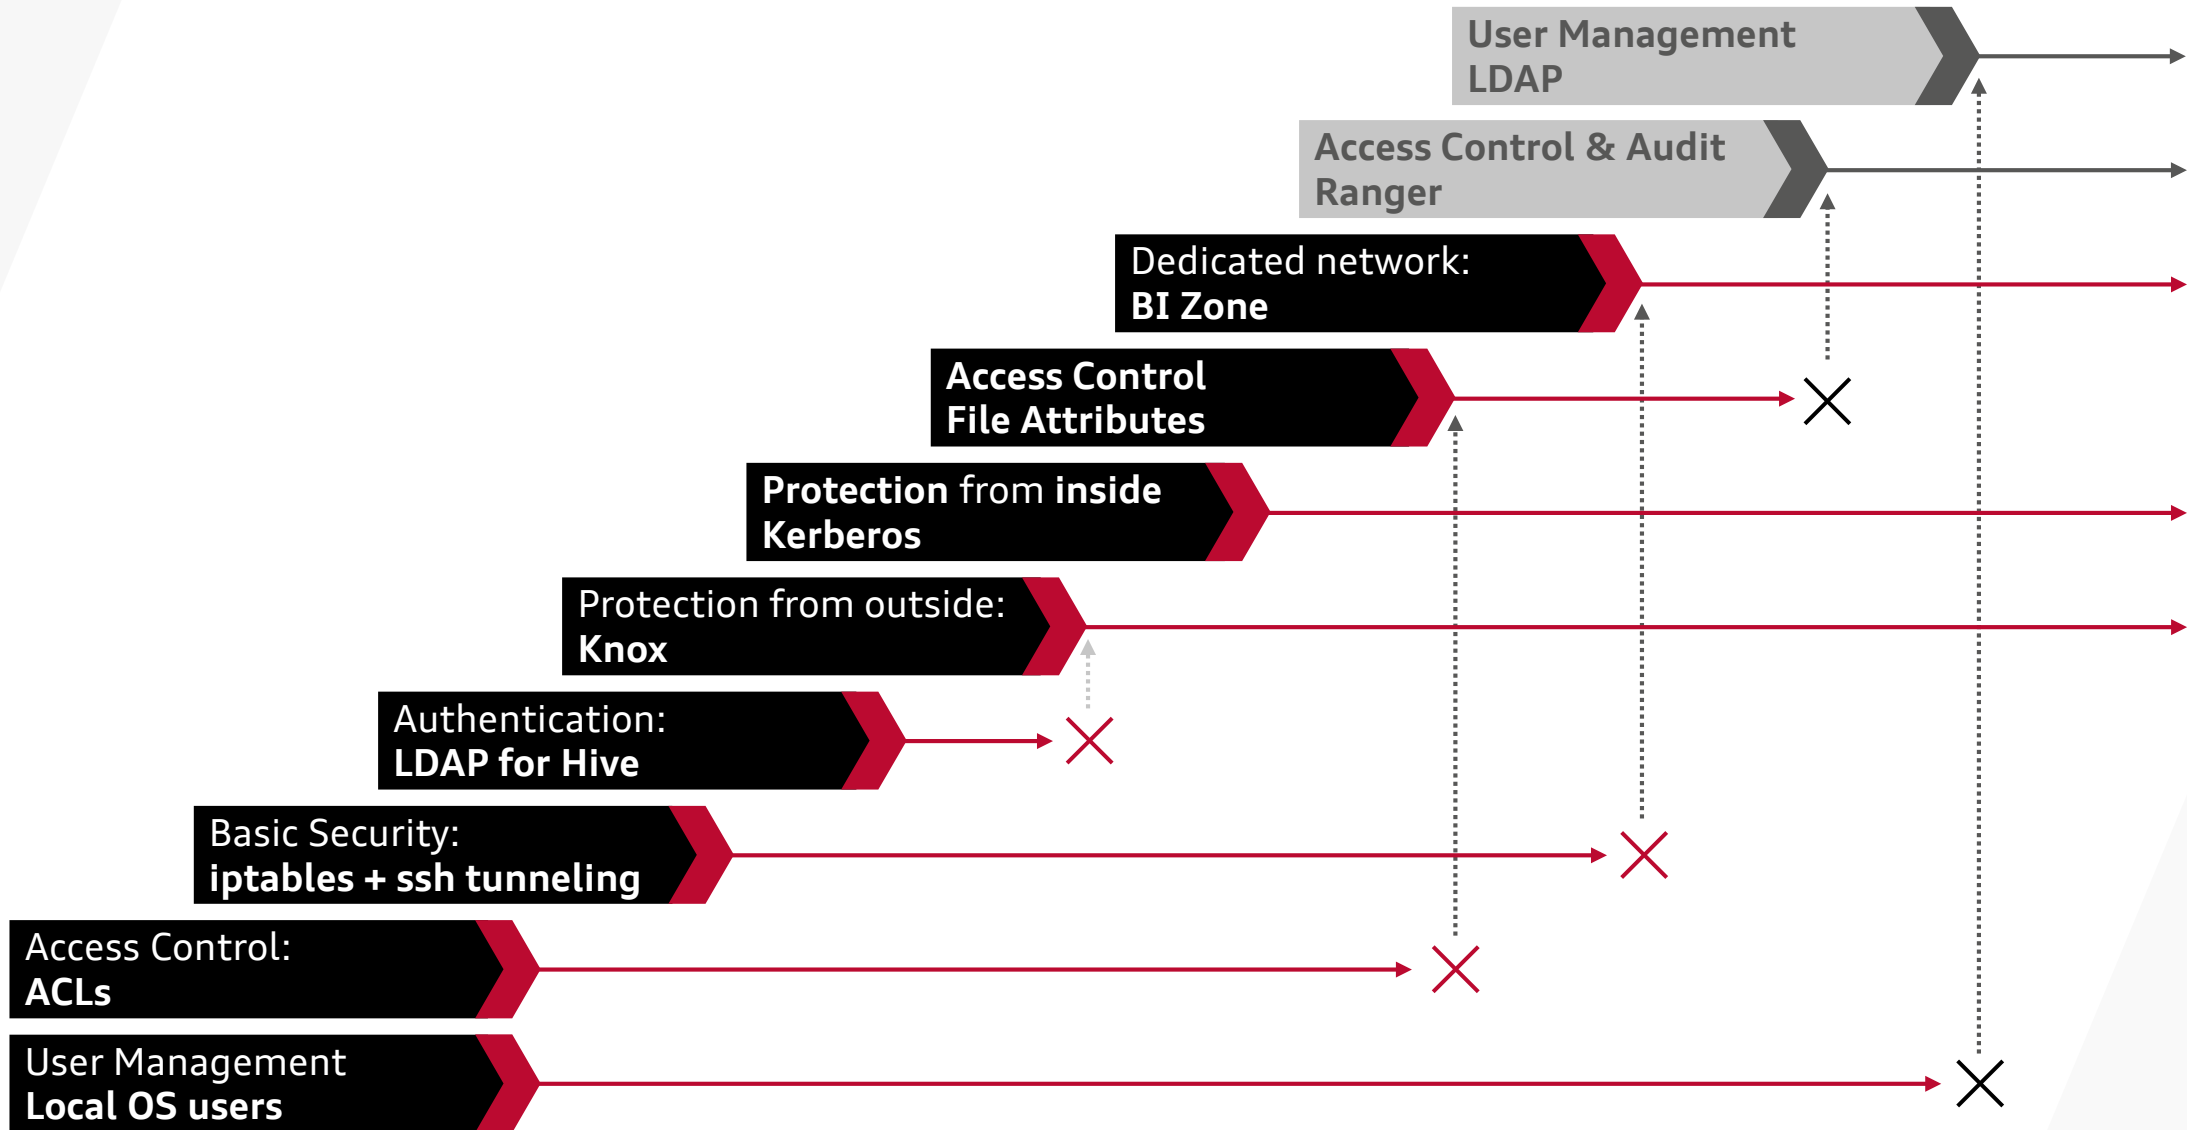
COMPANY CAR ANALYSIS



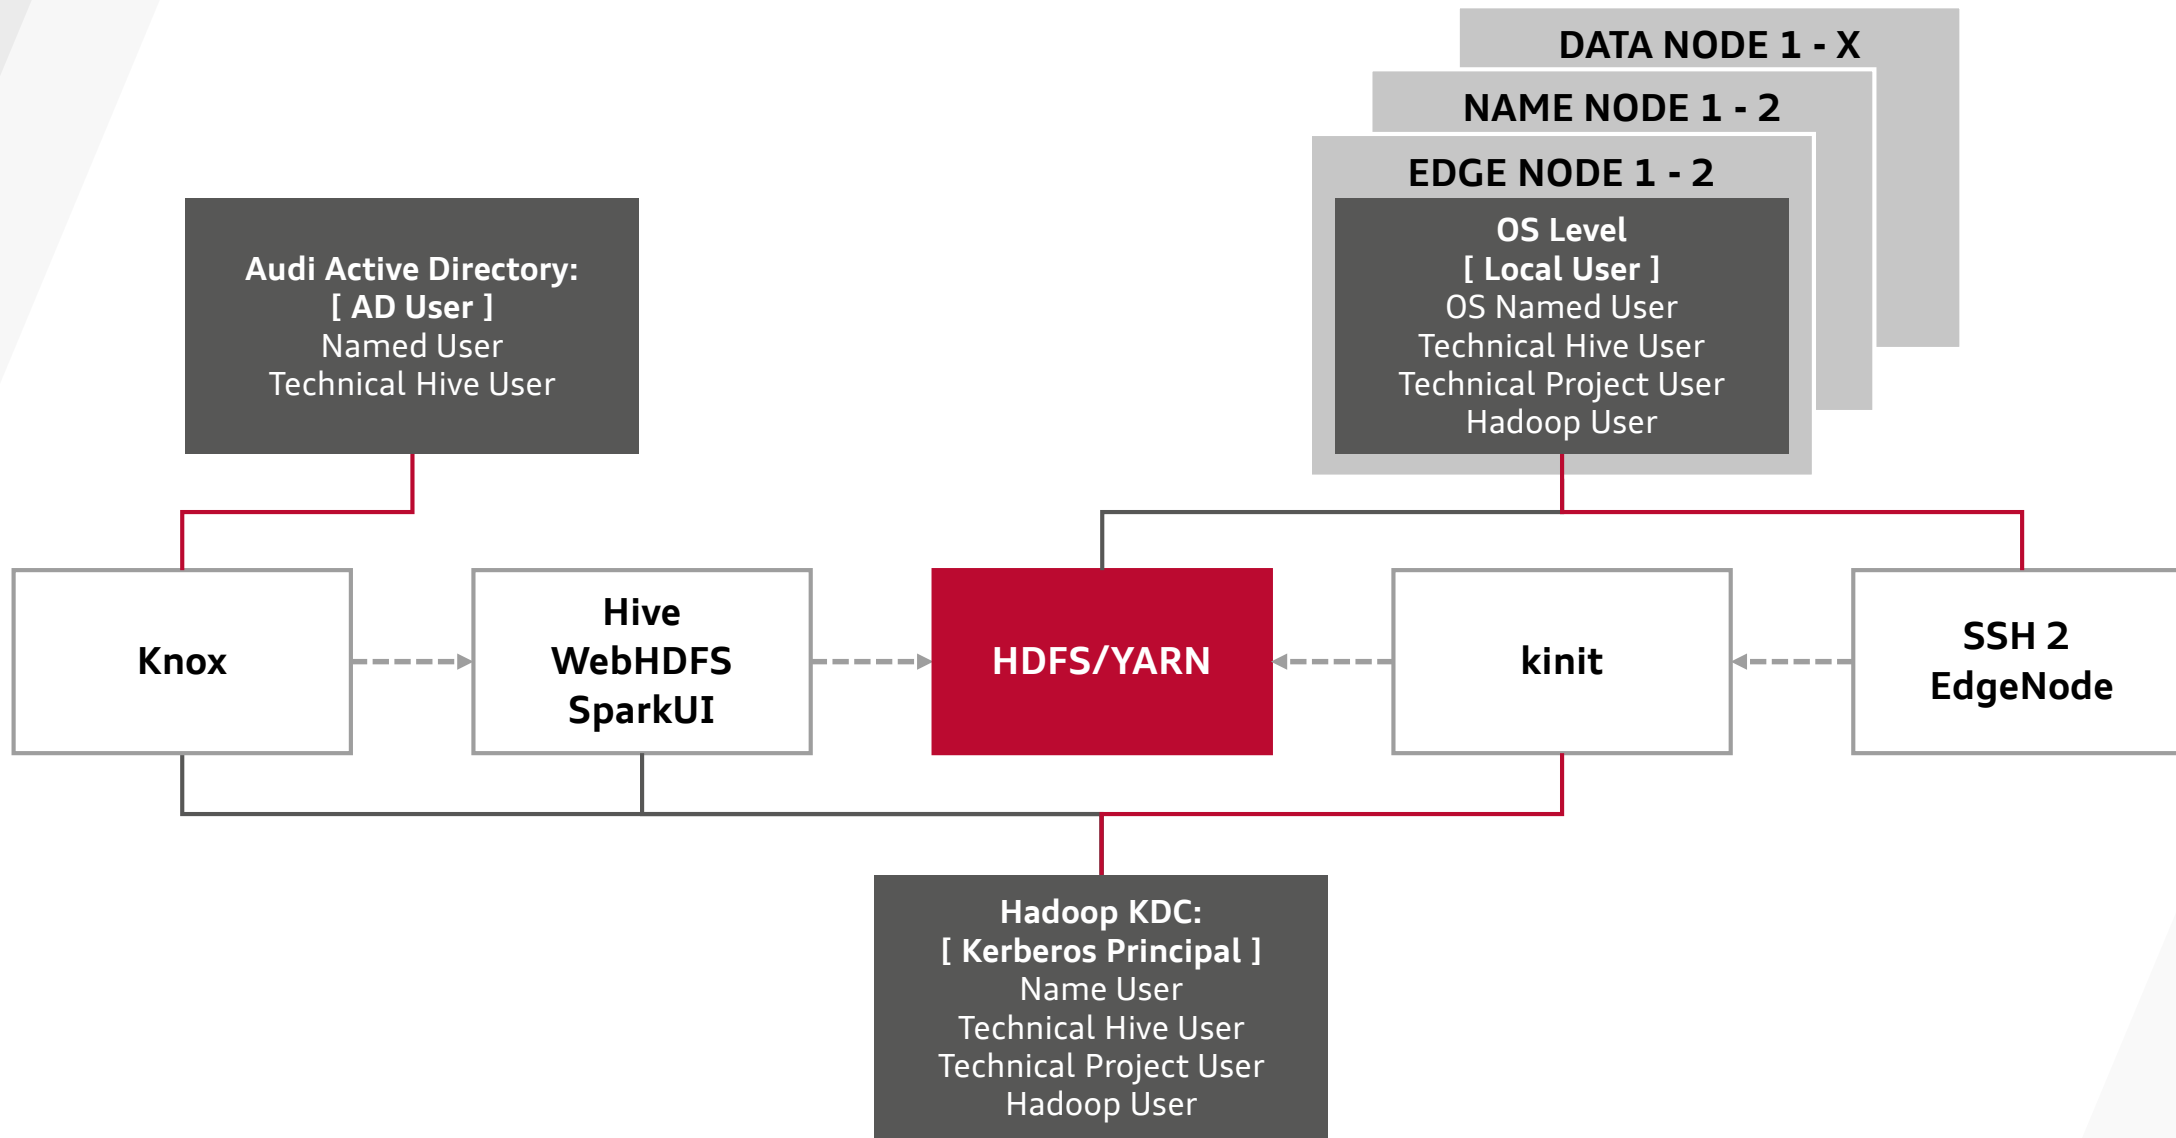
ENTERPRISE INTEGRATION VS SPEED OF DELIVERY

Securing the Cluster as multi-tenant environment

Step by step by step towards our target architecture ...



Password Hell



Legend:

— password required

— no password required

-----> next step

DATA INGESTION

Data ingestion: technical requirements from projects, security and ops



INGESTION

- » Streaming data
- » Batch data
- » easy writing to HDFS/DWH



DECOUPLING

- » Data Sources should not directly be coupled to analytical backend jobs
- » This allows adding new analytical jobs without changing the source



HA & BUFFERING

- » Data ingestion must be available 24x7
- » Data must be buffered (persisted) in case backend or backend job is not available



SECURITY

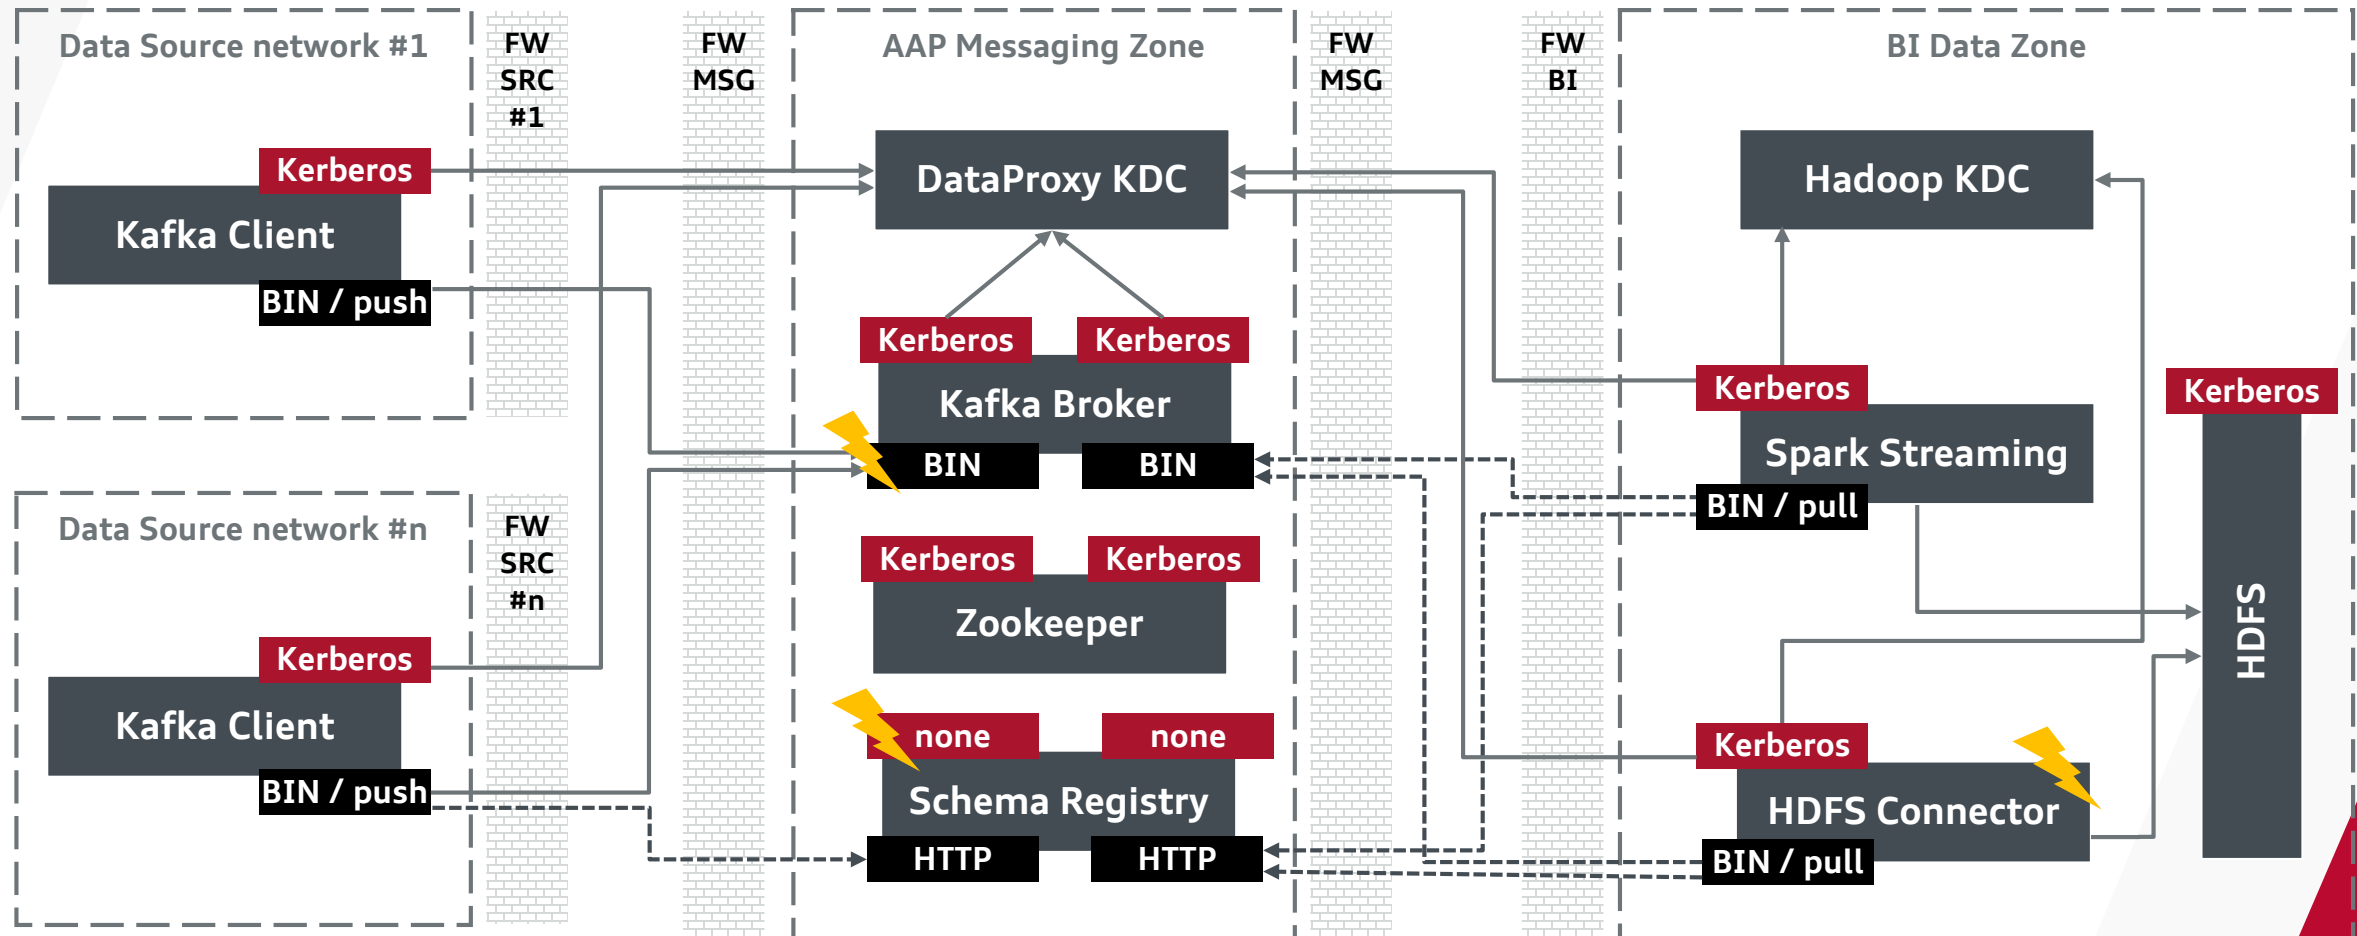
- » Source systems must not connect directly to the data zone (Hadoop, DWH) – by IT Sec
- » Authentication + Data in motion encryption (multi tenancy)
- » Protocol must be auditable
- » Some data sources run in the cloud



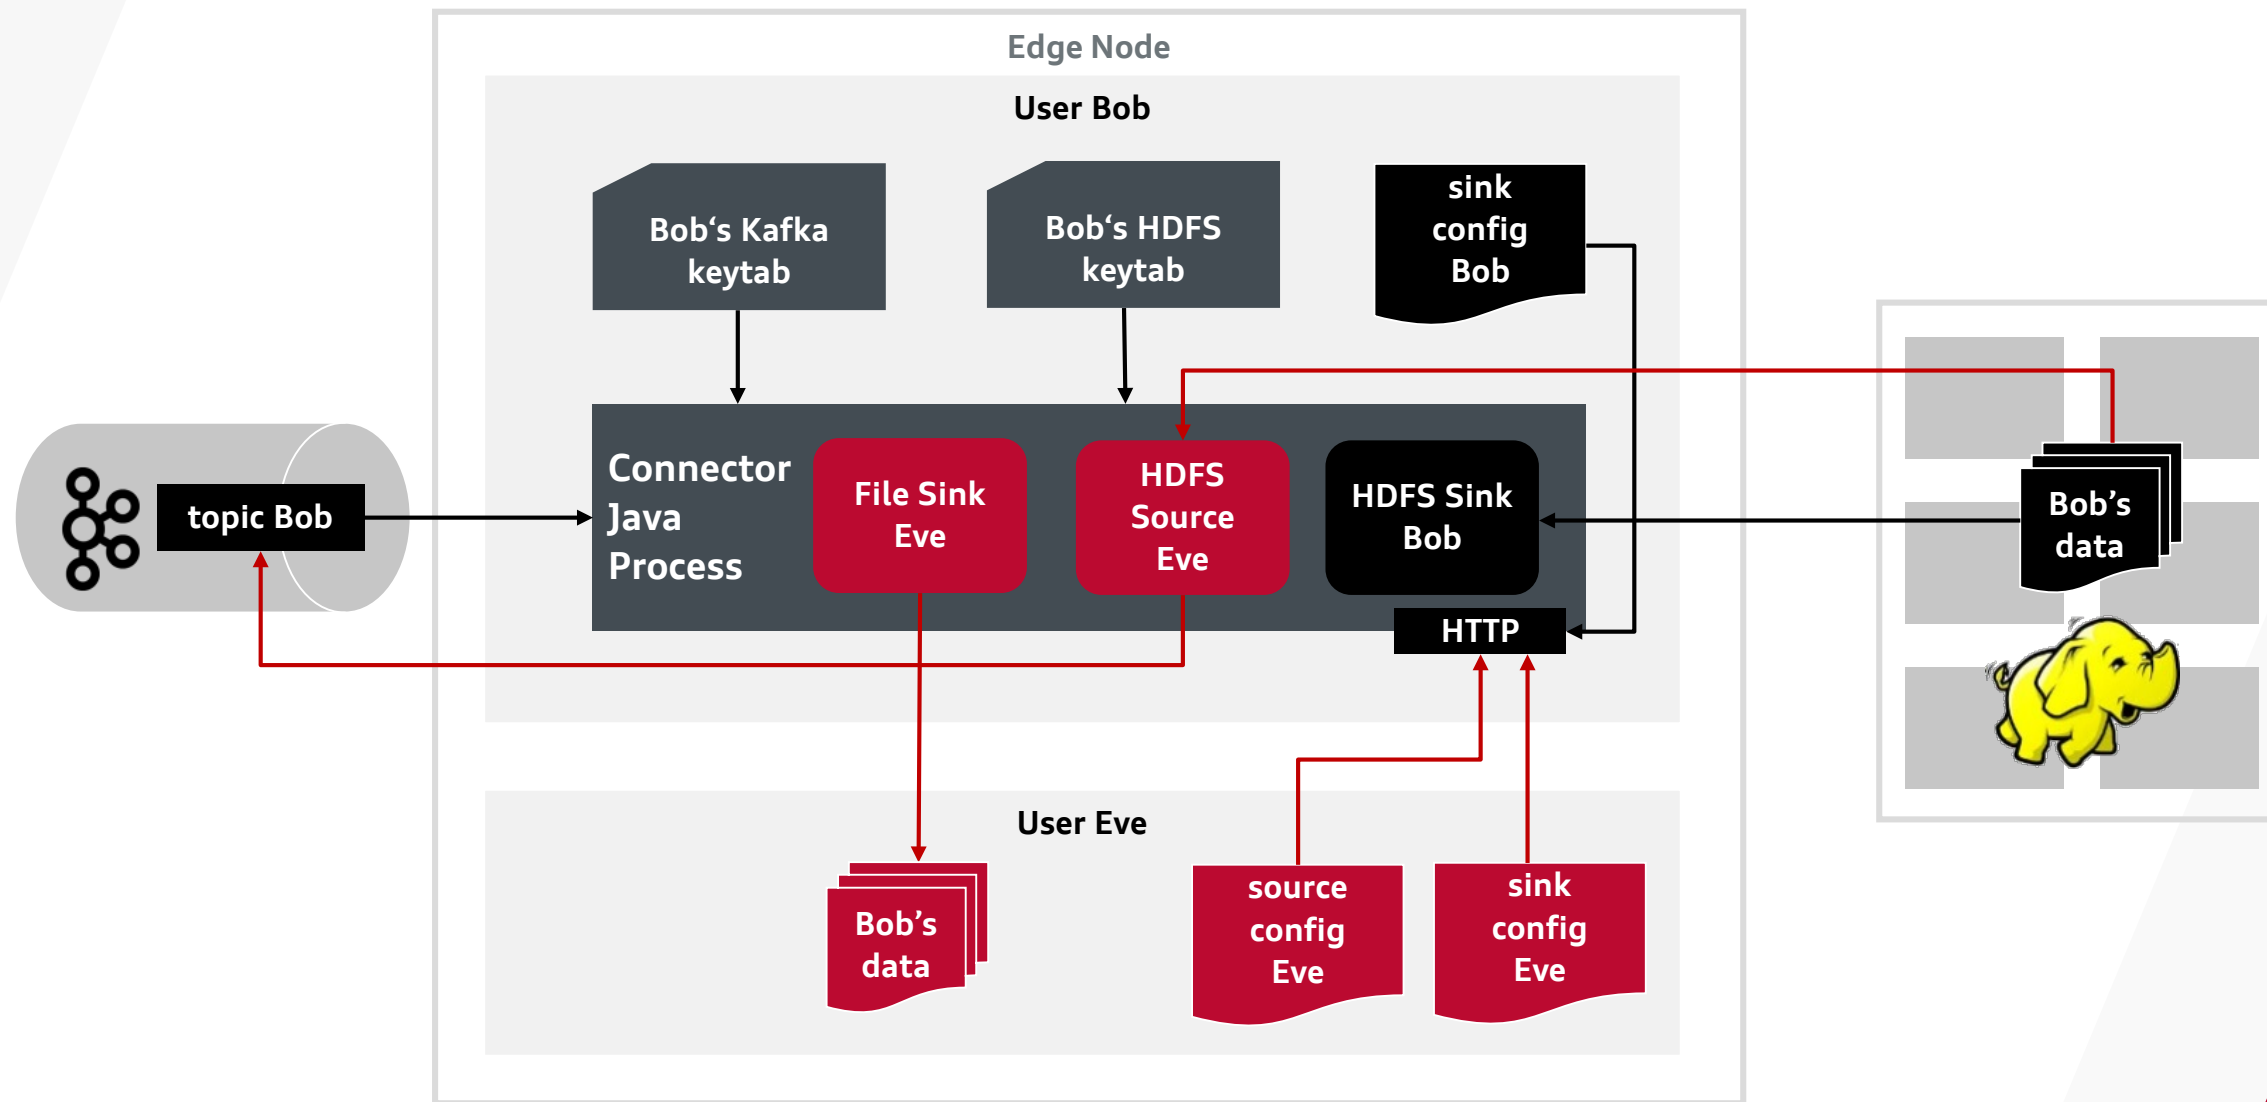
SCALABILITY

- » Amount of data will increase over time for most projects
- » Number of projects will increase

Solution: Kerberized Confluent Kafka Platform

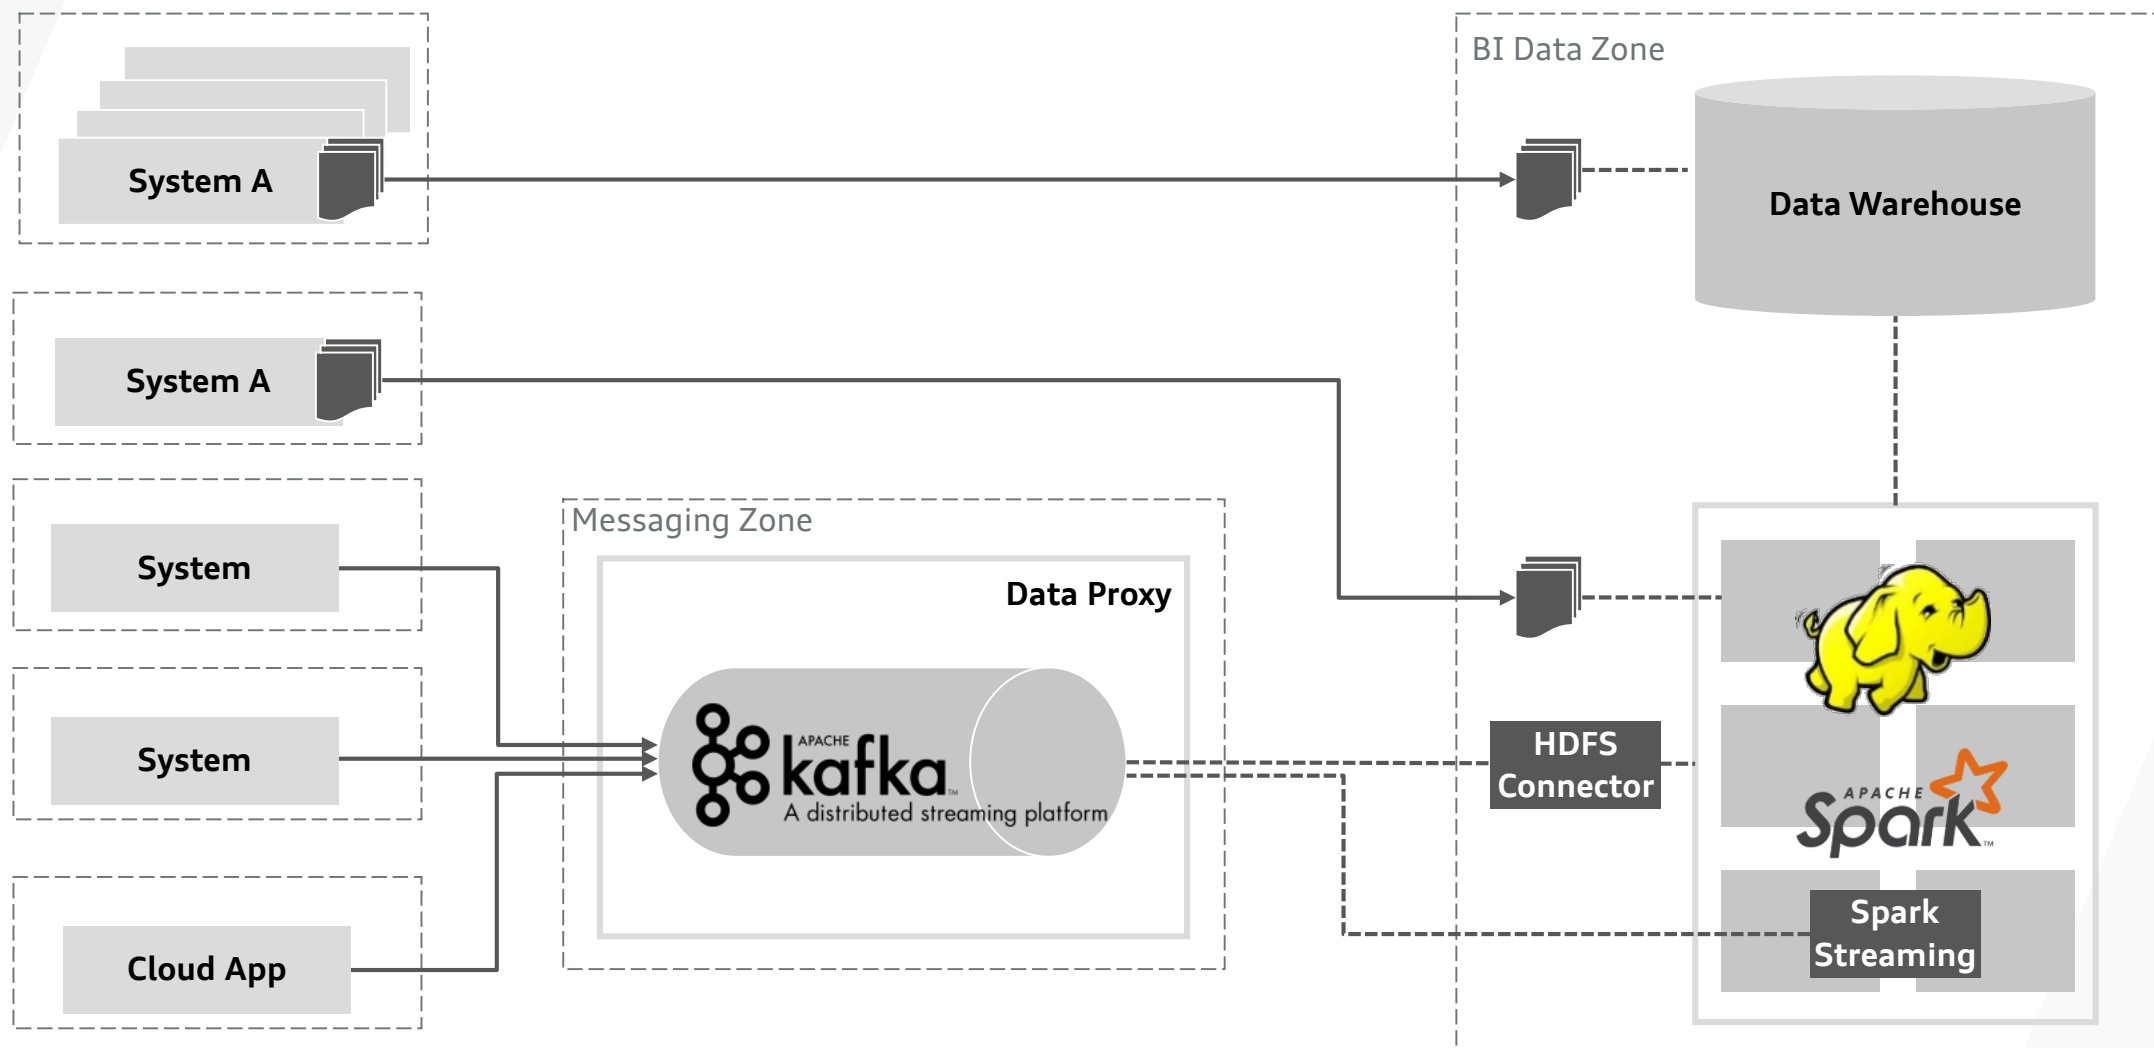


Kafka Distributed Connector: unsecured REST API

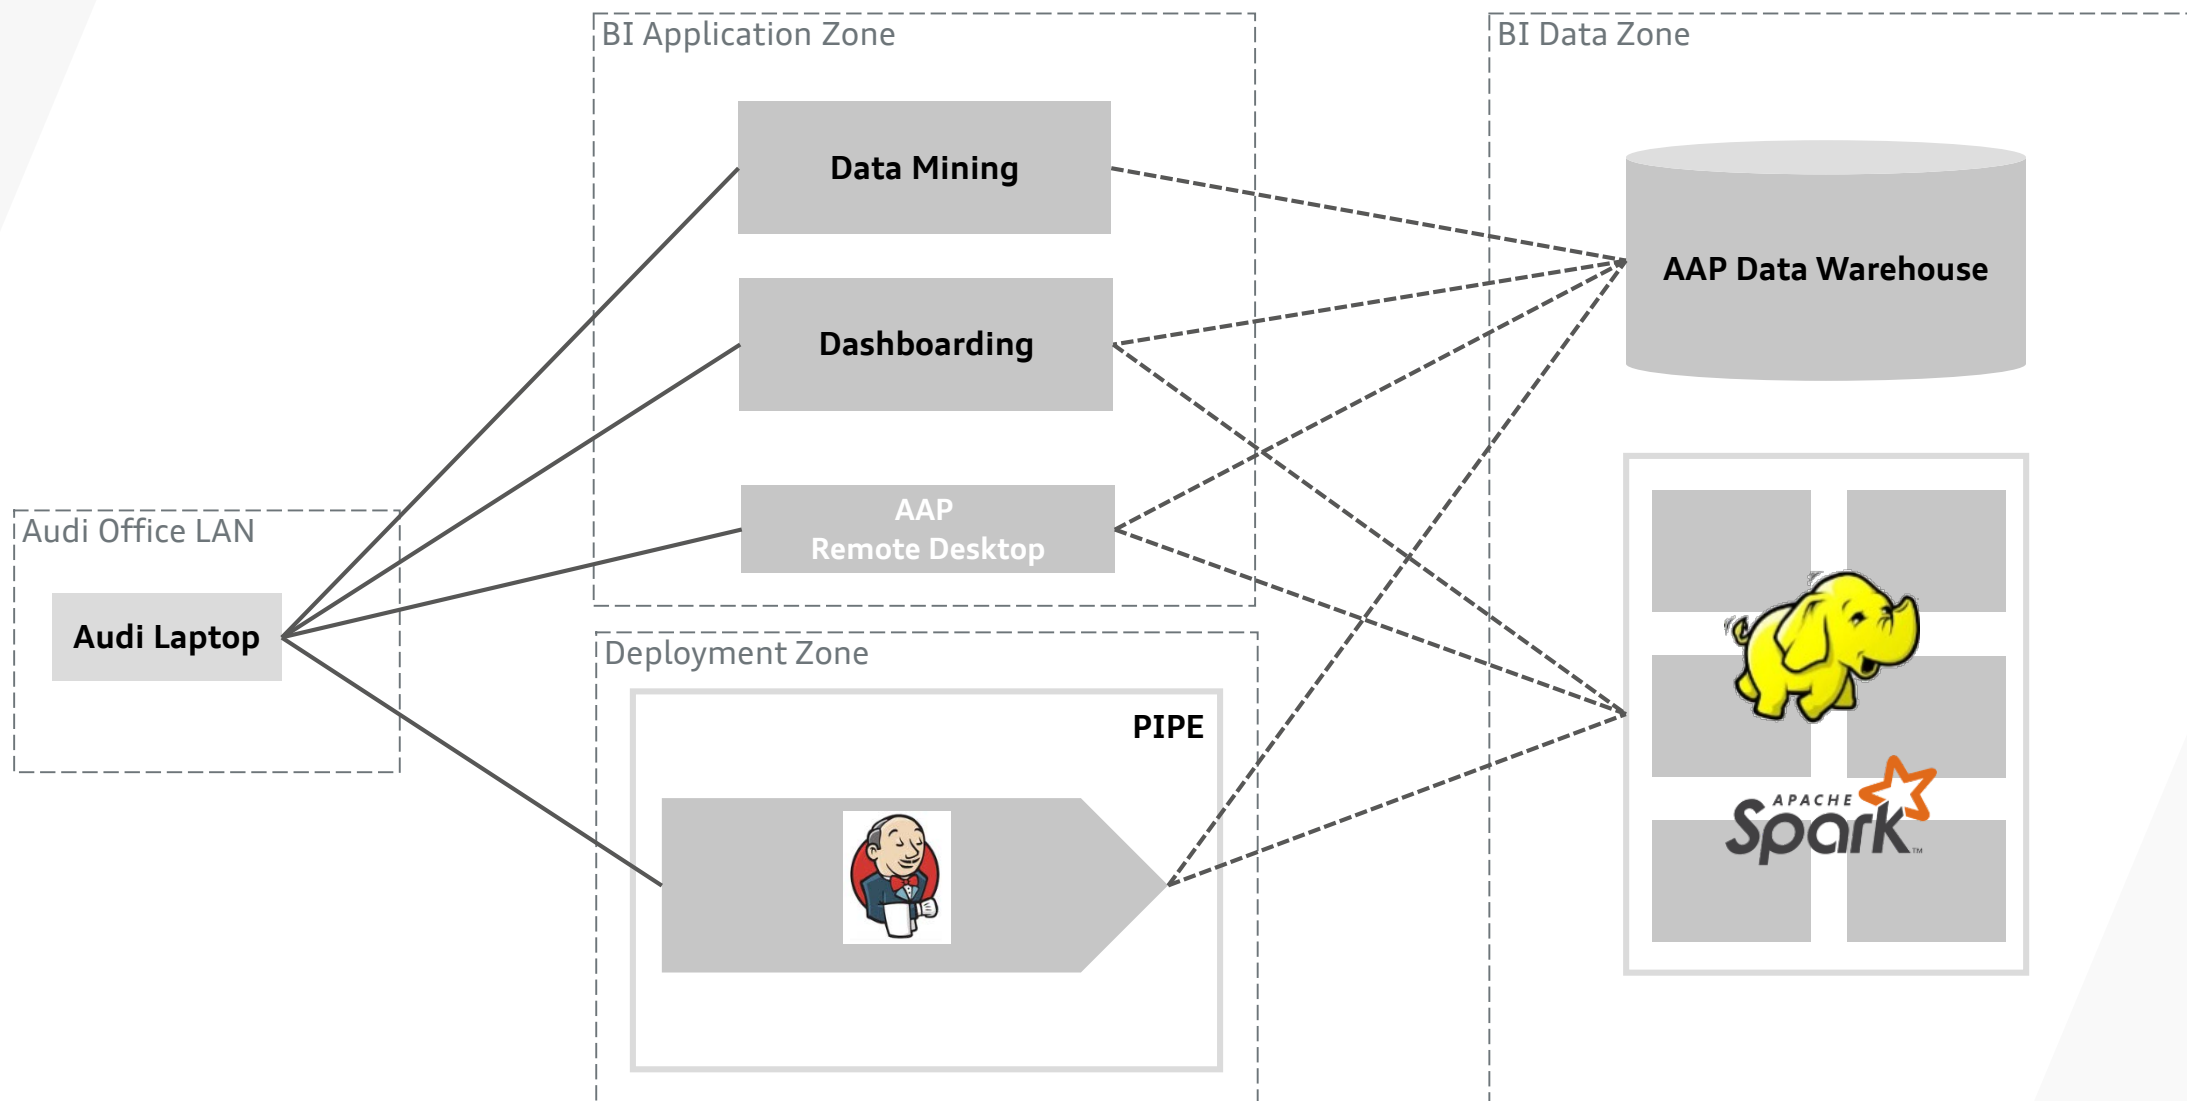


TODAY CURRENT STATE

Architecture & Network Zones – Data Ingestion



Architecture & Network Zones – User & Developer Access



Hadoop Cluster Sizing Production 2017



PROD

Hadoop	per node	Sum
# data nodes	1	12
RAM	512 GB	6 TB
Cores	24	288
HDD*	96 TB	9.216 TB

PROD

Kafka	per node	Sum
# broker nodes	1	4
RAM	32 GB	128 GB
Cores	6	24
HDD*	4 TB	16 TB

* Raw Capacity without replication and FS overhead!

Organisational Tasks

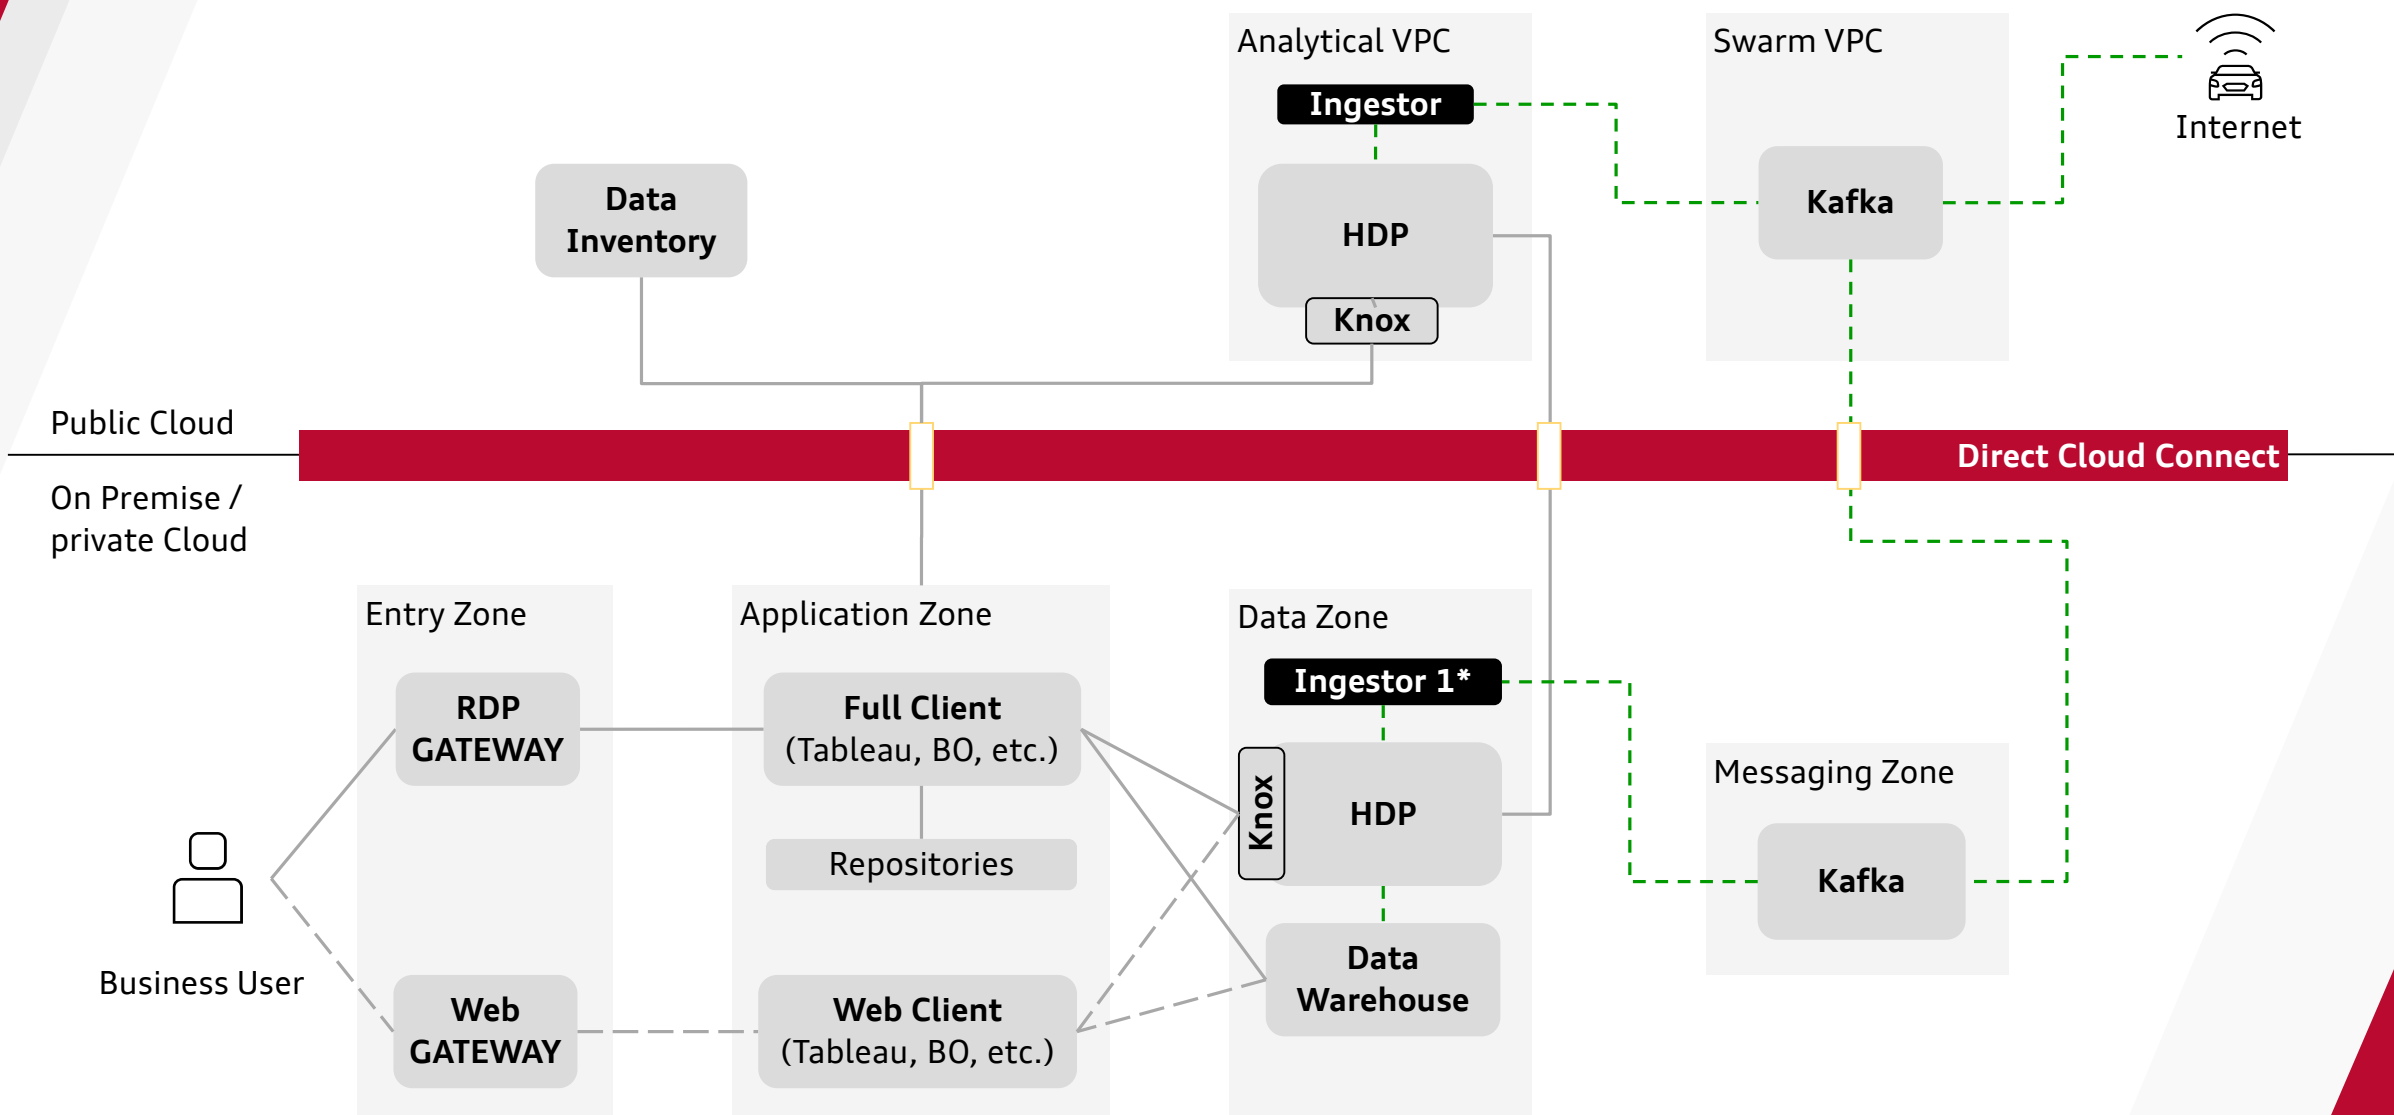
Current state

Organisational Tasks

- Data Ownership & Data Governance
(Data Domain Modell with clear responsibility in each domain)
- Lifecycle Management for each Shared Service in strong collaboration with the projects and programs
- Defined SLAs for each Shared Service based on general availability, data loss, confidentiality and verifiability
- Different Development Lifecycle between car and backend systems
- Use of Open Source Software and Support requirements from IT continuity
- Balance between multi tenant environment and flexibility
- Very long lifecycle of cars > 10 years with various built in software versions

TOMORROW WHAT'S UP NEXT

Hybrid Approach for the AAP





WE ARE HIRING

<https://www.audi.com/corporate/de/karriere/einstieg-bei-audi.html>

<https://karriere.audibusinessinnovation.com/>