



Data Ingestion Engine

Design Overview for Big Data Meetup

Agenda

- Introduction
- Data Ingestion Engine
 - Why do this? Isn't data ingestion just ETL?
 - Overview of Engine design
 - Advantages of this approach
- Future enhancements
- Q/A & Discussion (remainder)

Who is this guy?

- Steven Frobish, Senior Consultant at Daugherty
- 17 years in Data and Analytics
 - 3 years doing big data on mainframe before it was called “Big Data”
 - 4 years Informatica support
 - 4 years Data Integration Architect
 - 6 years in the Hadoop space
 - Ingestion, Cataloging, and Provisioning
 - Operational Analytics
 - Machine Learning



Motivation

- Analysts – Just give me the raw data!
- Auditors – Where is your sensitive information? (including free text)
- Records Retention – How long is your data going to stick around?
- Project Managers – I need 1800 tables moved in the next 2 weeks!
- DBA's – Your schema is going to change (and you won't know when)

Motivation - contd

- What did I learn:
 - As “Big Iron” analyst
 - Disk I/O is painfully slow, but ‘parallelizable’
 - Joins are expensive
 - As Informatica Support Engineer:
 - Roughly 30% of mappings were ‘pass-thru’
 - Another 50% were similarly “doing the same thing over and over”
 - ...but on slightly different data
 - As Data Integration Architect:
 - Most consumers can make sense of data on their own
 - They just want everything, all the time
 - Even highly modeled stores have common features that can be generalized

Motivation - contd

- What did I learn:
 - As Data Lake Builder
 - Lines of Business almost always need change data if available
 - They are happy with schema on read
 - but don't break their stuff
 - They will inevitably have their favorite format
 - which they will discard immediately in favor of consistency
 - Audits happen
 - Security will always tell you that you are doing it wrong
 - unless you can plug in any level of encryption they can think of...
 - (then literally anything other than plaintext will do)

Why not ETL?

- Connecting to data is the easy part
 - JDBC
 - SFTP
 - REST API's
 - Loading to Hadoop from edge node
- ETL tool's main product is 'T' (no 'T' = no ETL)
- 1 table = 1 mapping
- Schema definition is part and parcel to the mapping
- Resource based Scheduling is difficult (impossible?)



Why not Spark?

- Spark is a parallel processing framework
 - It is not free performance or instant throughput
 - “Accessing the power of the cluster” is not needed here
- Problem of ingesting data is inherently not parallel
 - Read a single source of data and write to a single target
 - Any parallelism added in the middle is only adding complexity, not improving throughput

What about NiFi?

- NiFi is cool
 - Has many features Engine does not explore
 - Back-pressure, fine-grained access, write-ahead logs
 - An event driven architecture
 - Typically requires more “plumbing” to set up
- Flows are still 1:1 to sources and actions are not metadata driven
- Costly to setup many flows at once

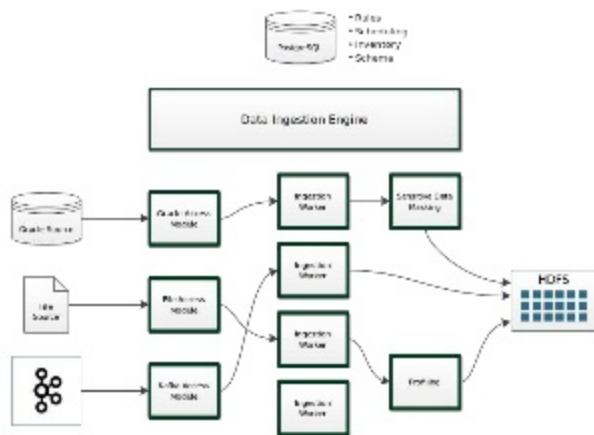


What about SQOOP?

- Actually SQOOP is in the current Engine as an Access Module
- It has many features we want:
 - Incremental extracts
 - Captures schema information
- And several shortcomings:
 - No resource based scheduling
 - Doesn't track inventory
 - Field level masking is very difficult

The Data Ingestion Engine

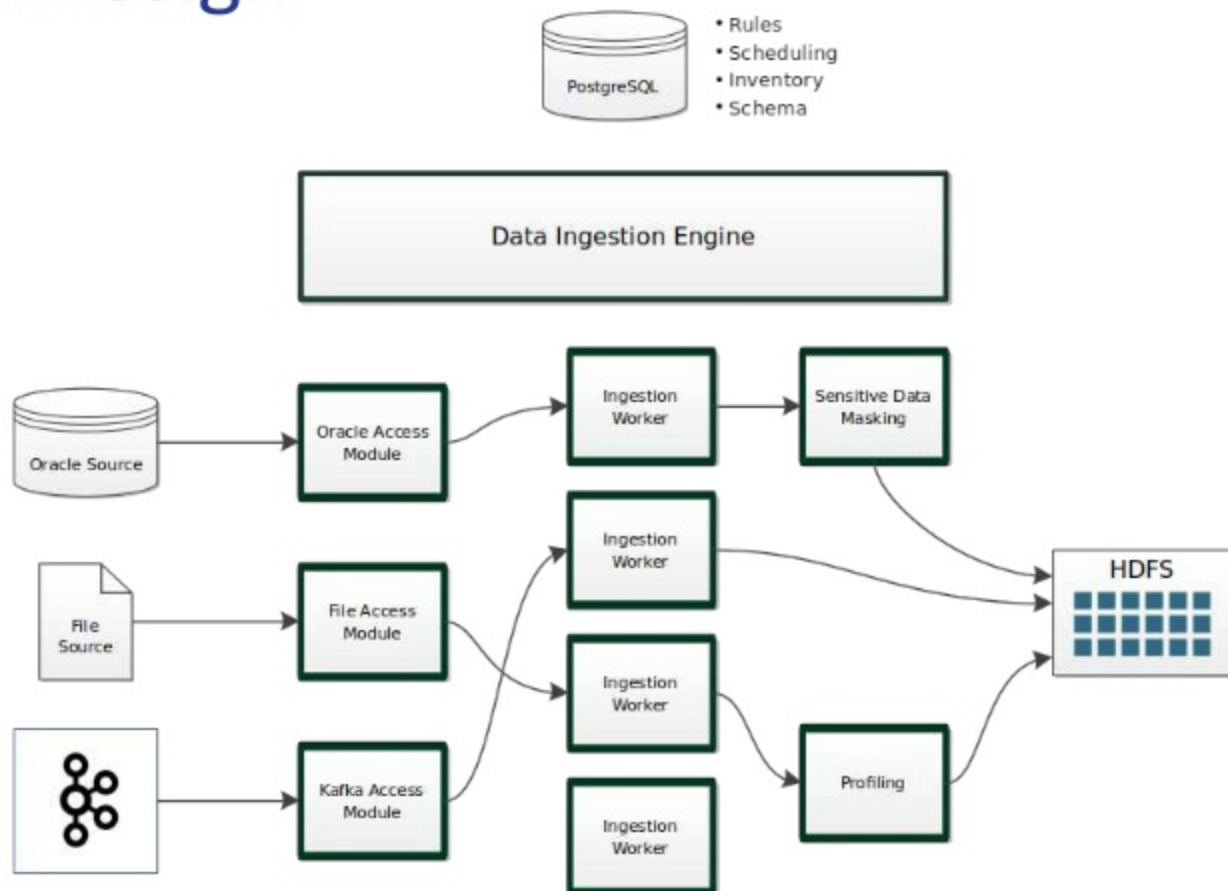
- Is really just a driver program that:
 - Queries a repository
 - Decides to run jobs
 - Time based
 - ...but conscious of existing work
 - Launches worker threads



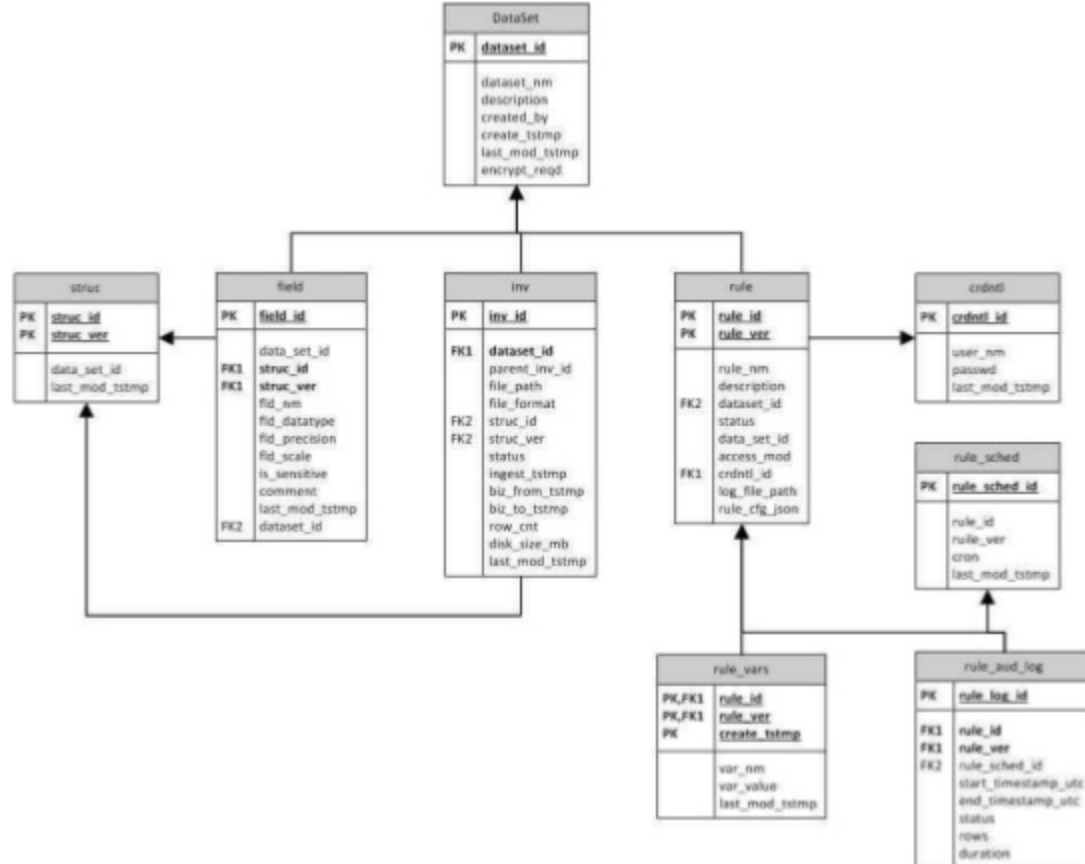
Workers

- Find work to do from the queue
- Load up appropriate access module
- Track:
 - Logging
 - Runtime metadata
 - Inventory metadata
- Occasionally:
 - Transform sensitive data
 - Capture profile stats
- Report errors
- Find *more* work to do from queue

Engine Design



Metadata Repository



Access Modules

- The code necessary to access a given source
 - Oracle
 - SFTP
 - Kafka
- Each module is responsible for:
 - Pulling data
 - Acquiring schema information
 - Recording basic statistics - row/record count, byte size, time taken

Transform Modules

- Use schema information to:
 - Mask data
 - Transform data to a given format
 - Optionally capture profile stats
- Did someone mention sensitive data?
 - HIPAA, PCI, PII, etc.
 - Would some column tags help?
 - Explicitly tag column name (e.g. 'colA' tagged as "is_sensitive")
 - Regex on column name (e.g. like '.*phone.*', or 'b(irth)?_?date')
 - Have some regexes for free text
 - Regex on content (e.g. '\d{3}-?\d{2}-?\d{4}')

Metadata Driven Approach Benefits

- Generate many rules quickly
- Imagine an Access Module for which the output is more rules metadata
 - Scan Oracle Catalog for a given schema
 - Generate rules/datasets for every table
 - Let the access module do the heavy lifting
 - You just need to tell it where to go and provide credentials

Metadata Driven Approach Benefits

- Capture schema changes immediately
- Use inventory data to trigger consumer processes
- Full lineage & Audit logs accessible in a single place
- Resource based scheduling and connection throttling
- Information Retention driven by metadata instead of scanning data
- ... many more that become apparent only after you have built it

Some other stuff we are working on...

- Using KSQL to enrich data from Kafka
 - Engine will start (if necessary) KSQL server
 - With metadata generated KSQL***
- Provision data to an outside location***
- Launch Oozie workflow

*** Future feature

Questions?

Contact

steven.frobish@daugherty.com

Linked In