

AWS  
re:Invent

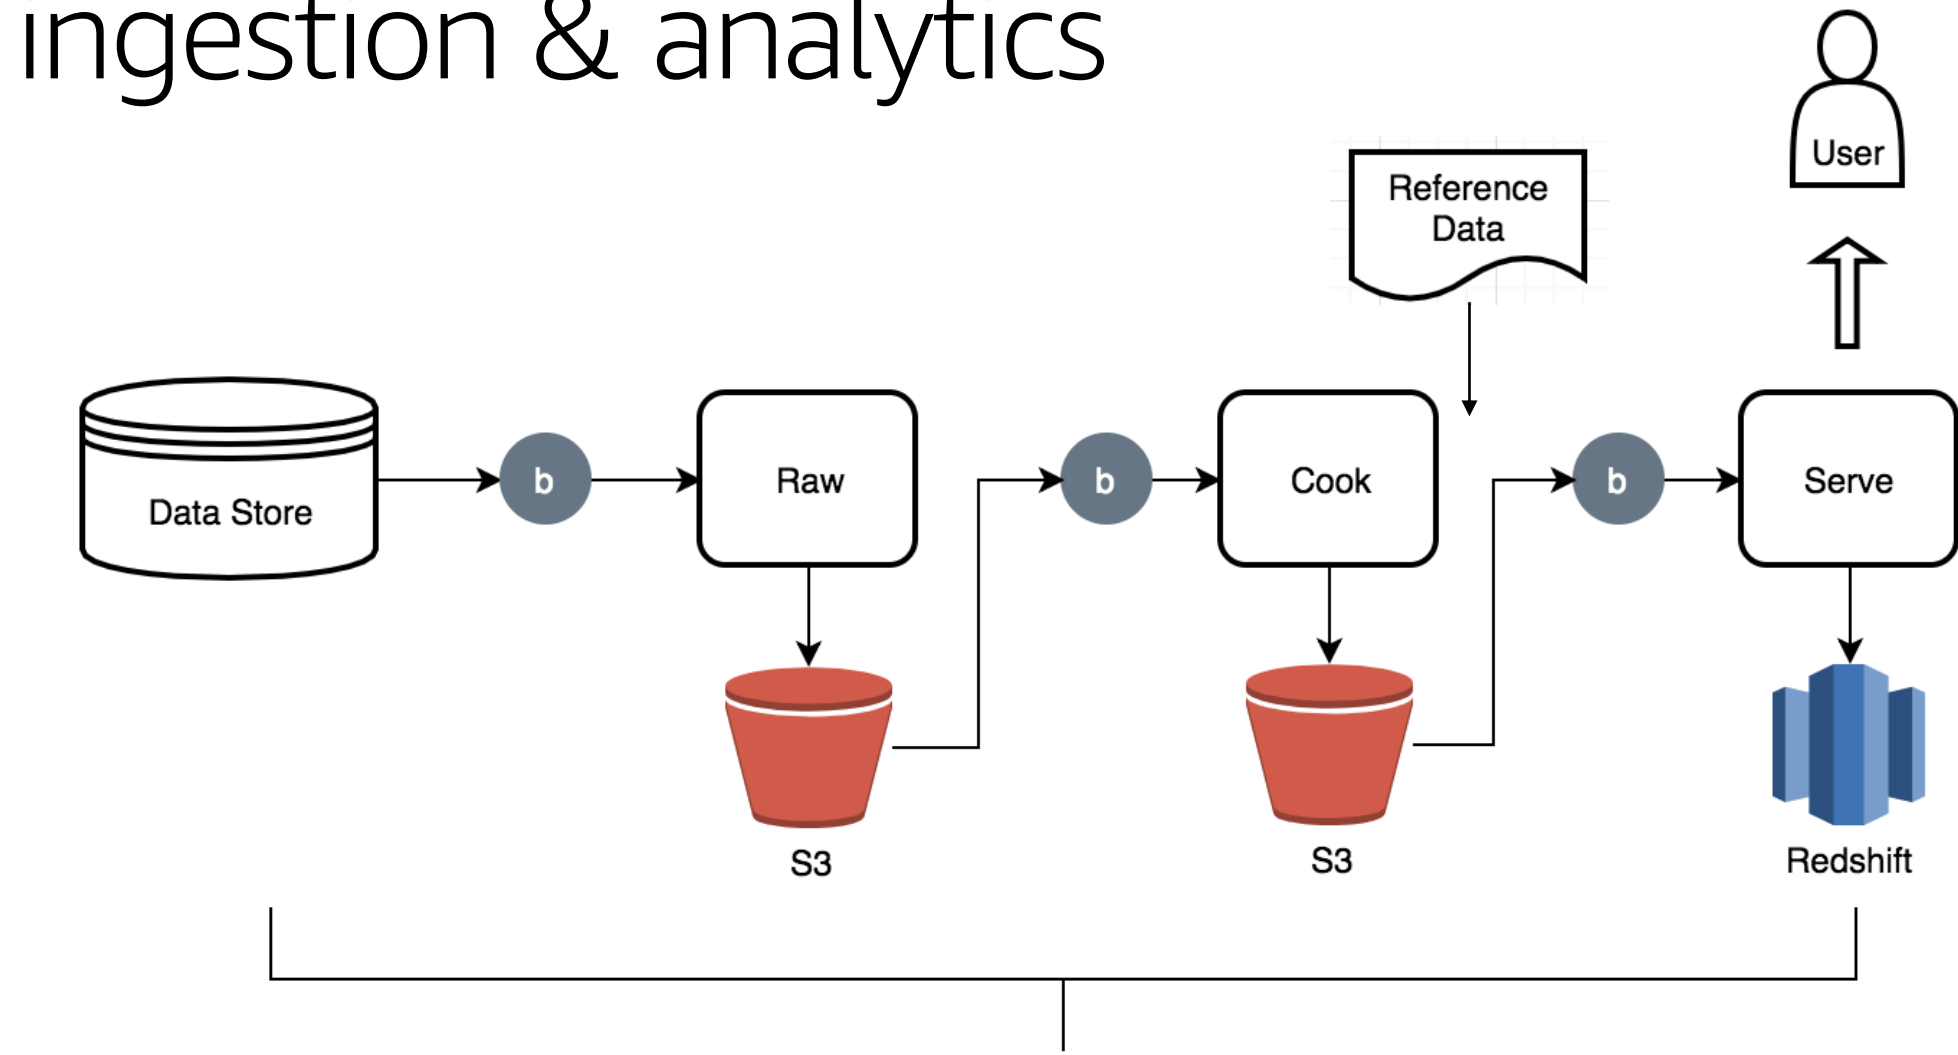
GP S C T 4 0 2

# Keeping the Pace with Data Ingestion

Paul Sears  
Partner Solutions Architect  
Amazon Web Services

Rohit Pujari  
Partner Solutions Architect  
Amazon Web Services

# Batch ingestion & analytics



**b** Batch Jobs

- Several hours
- Large batches

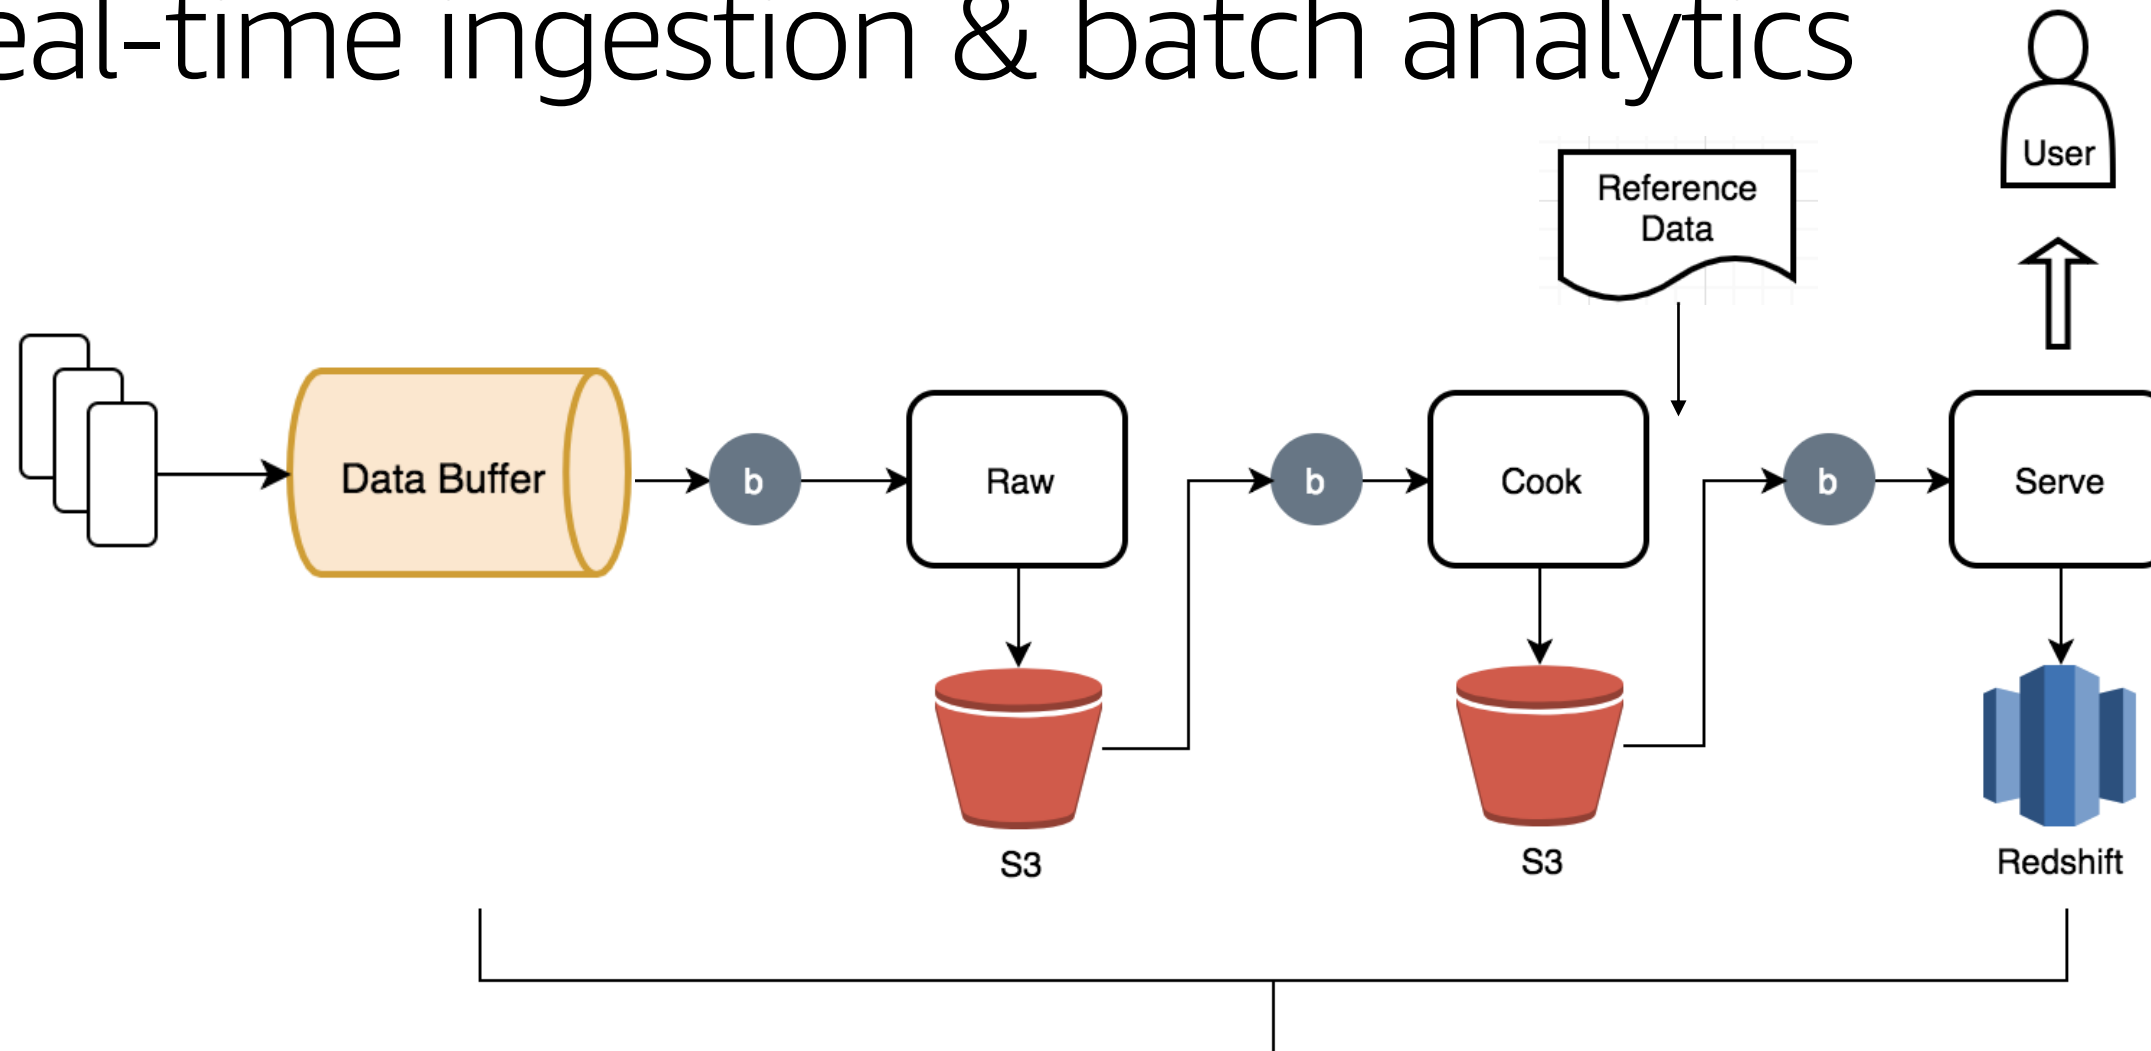


EMR



Glue ETL

# Real-time ingestion & batch analytics



**b** Batch Jobs

- Several hours
- Large batches
- Few hours
- Mini batches



EMR



Glue ETL

*value = data*

$$value = \frac{data}{time}$$

# Decision windows are shrinking

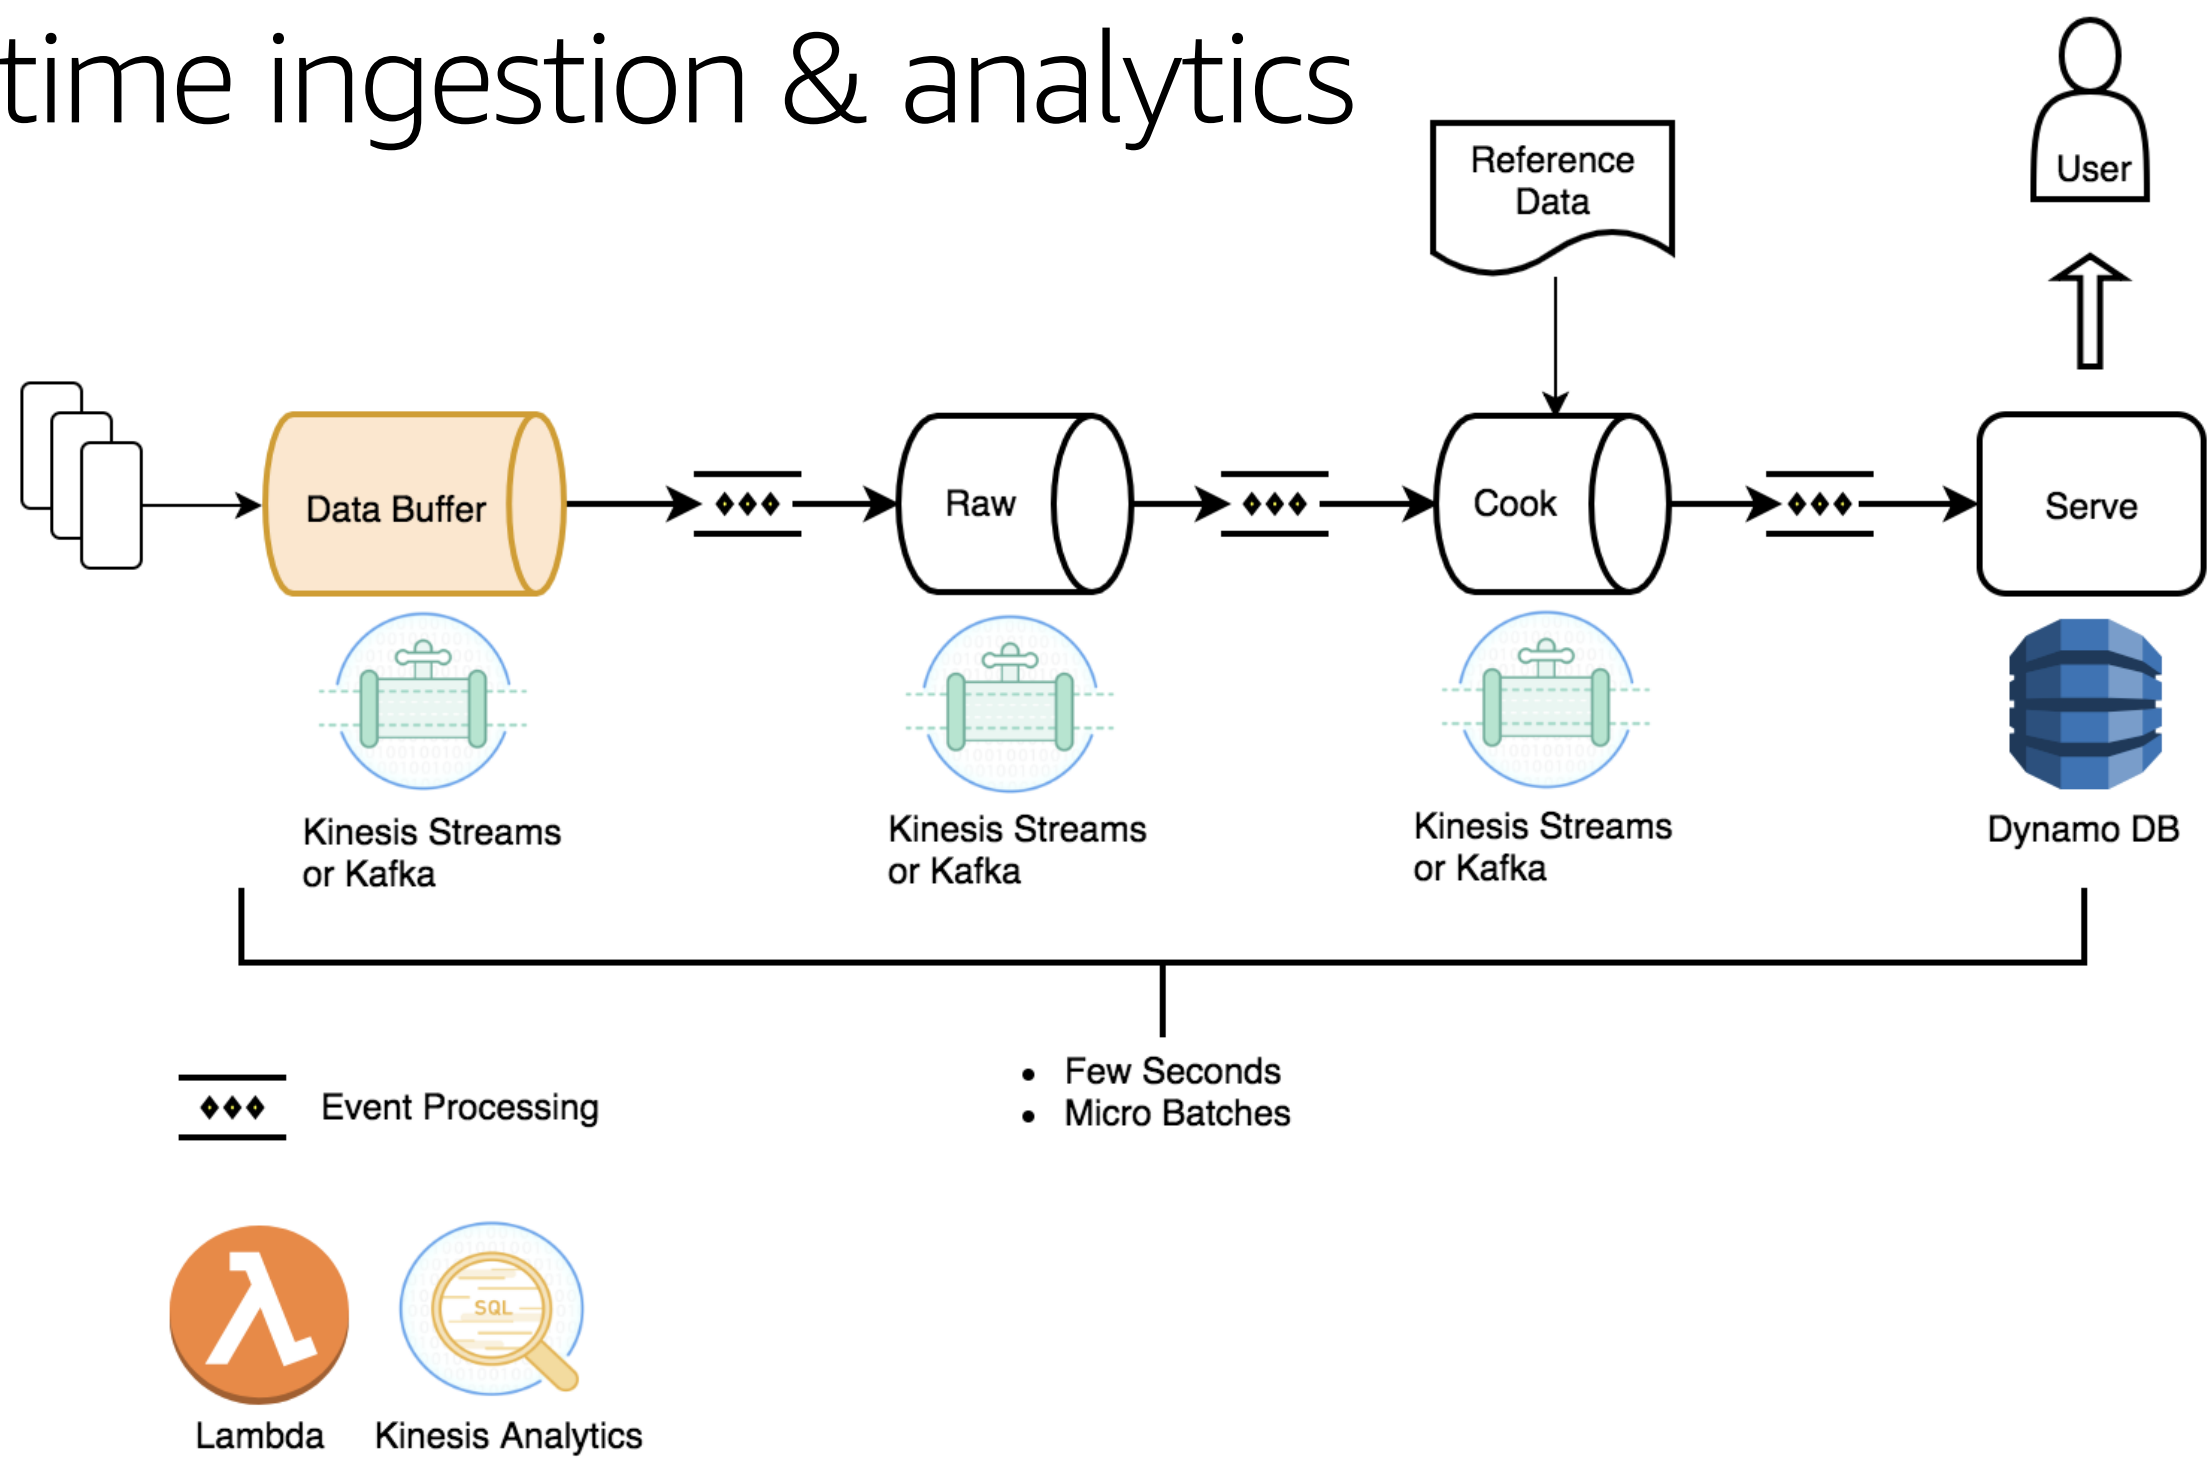
Hourly server logs → Real-time metrics

Weekly or monthly bills → Real-time spending alerts

Daily website clickstream → Real-time clickstream analysis


Daily fraud reports → Real-time detection

# Real-time ingestion & analytics





# Amazon Kinesis




## Amazon Kinesis Data Streams

Build custom applications  
that process and analyze  
streaming data



## Amazon Kinesis Data Analytics

Easily process and  
analyze streaming data  
with standard SQL



## Amazon Kinesis Data Firehose

Easily load streaming  
data into AWS

# Simple pattern for streaming data

## Data producer

Continuously creates data

Continuously writes data to a stream

Can be almost anything

## Streaming service

Durably stores data

Provides temporary buffer that preps data

Supports very high-throughput

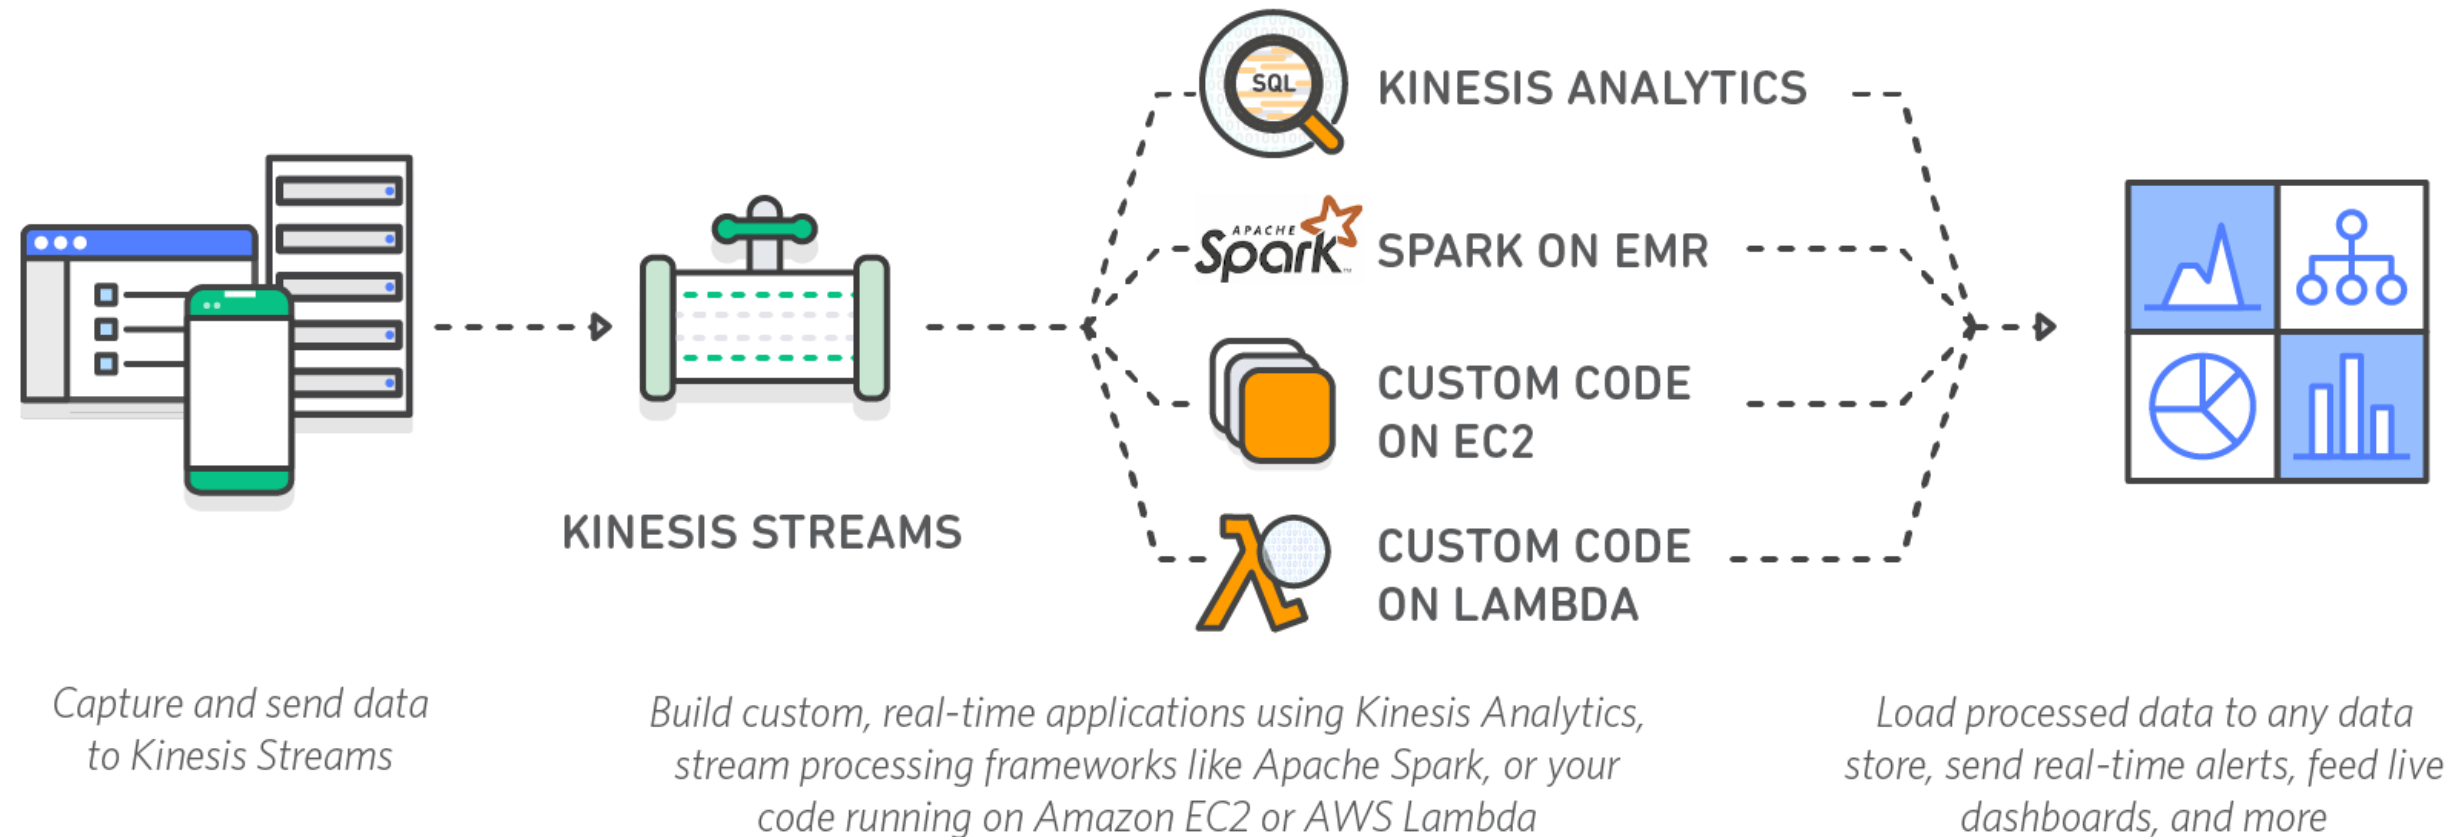
## Data consumer

Continuously processes data

Cleans, prepares, & aggregates

Transforms data to information

# Amazon Kinesis Data Streams



- Easy administration and low cost
- Build real-time applications with framework of choice
- Secure, durable storage

# Amazon Kinesis Data Analytics




*Capture streaming data with  
Kinesis Streams or Kinesis Firehose*

*Run standard SQL queries  
against data streams*

*Kinesis Analytics can send processed data to  
analytics tools so you can create alerts and  
respond in real-time*

- Powerful real-time applications
- Easy to use, fully managed
- Automatic elasticity

# Amazon Kinesis Data Firehose



*Capture and send data  
to Kinesis Firehose*

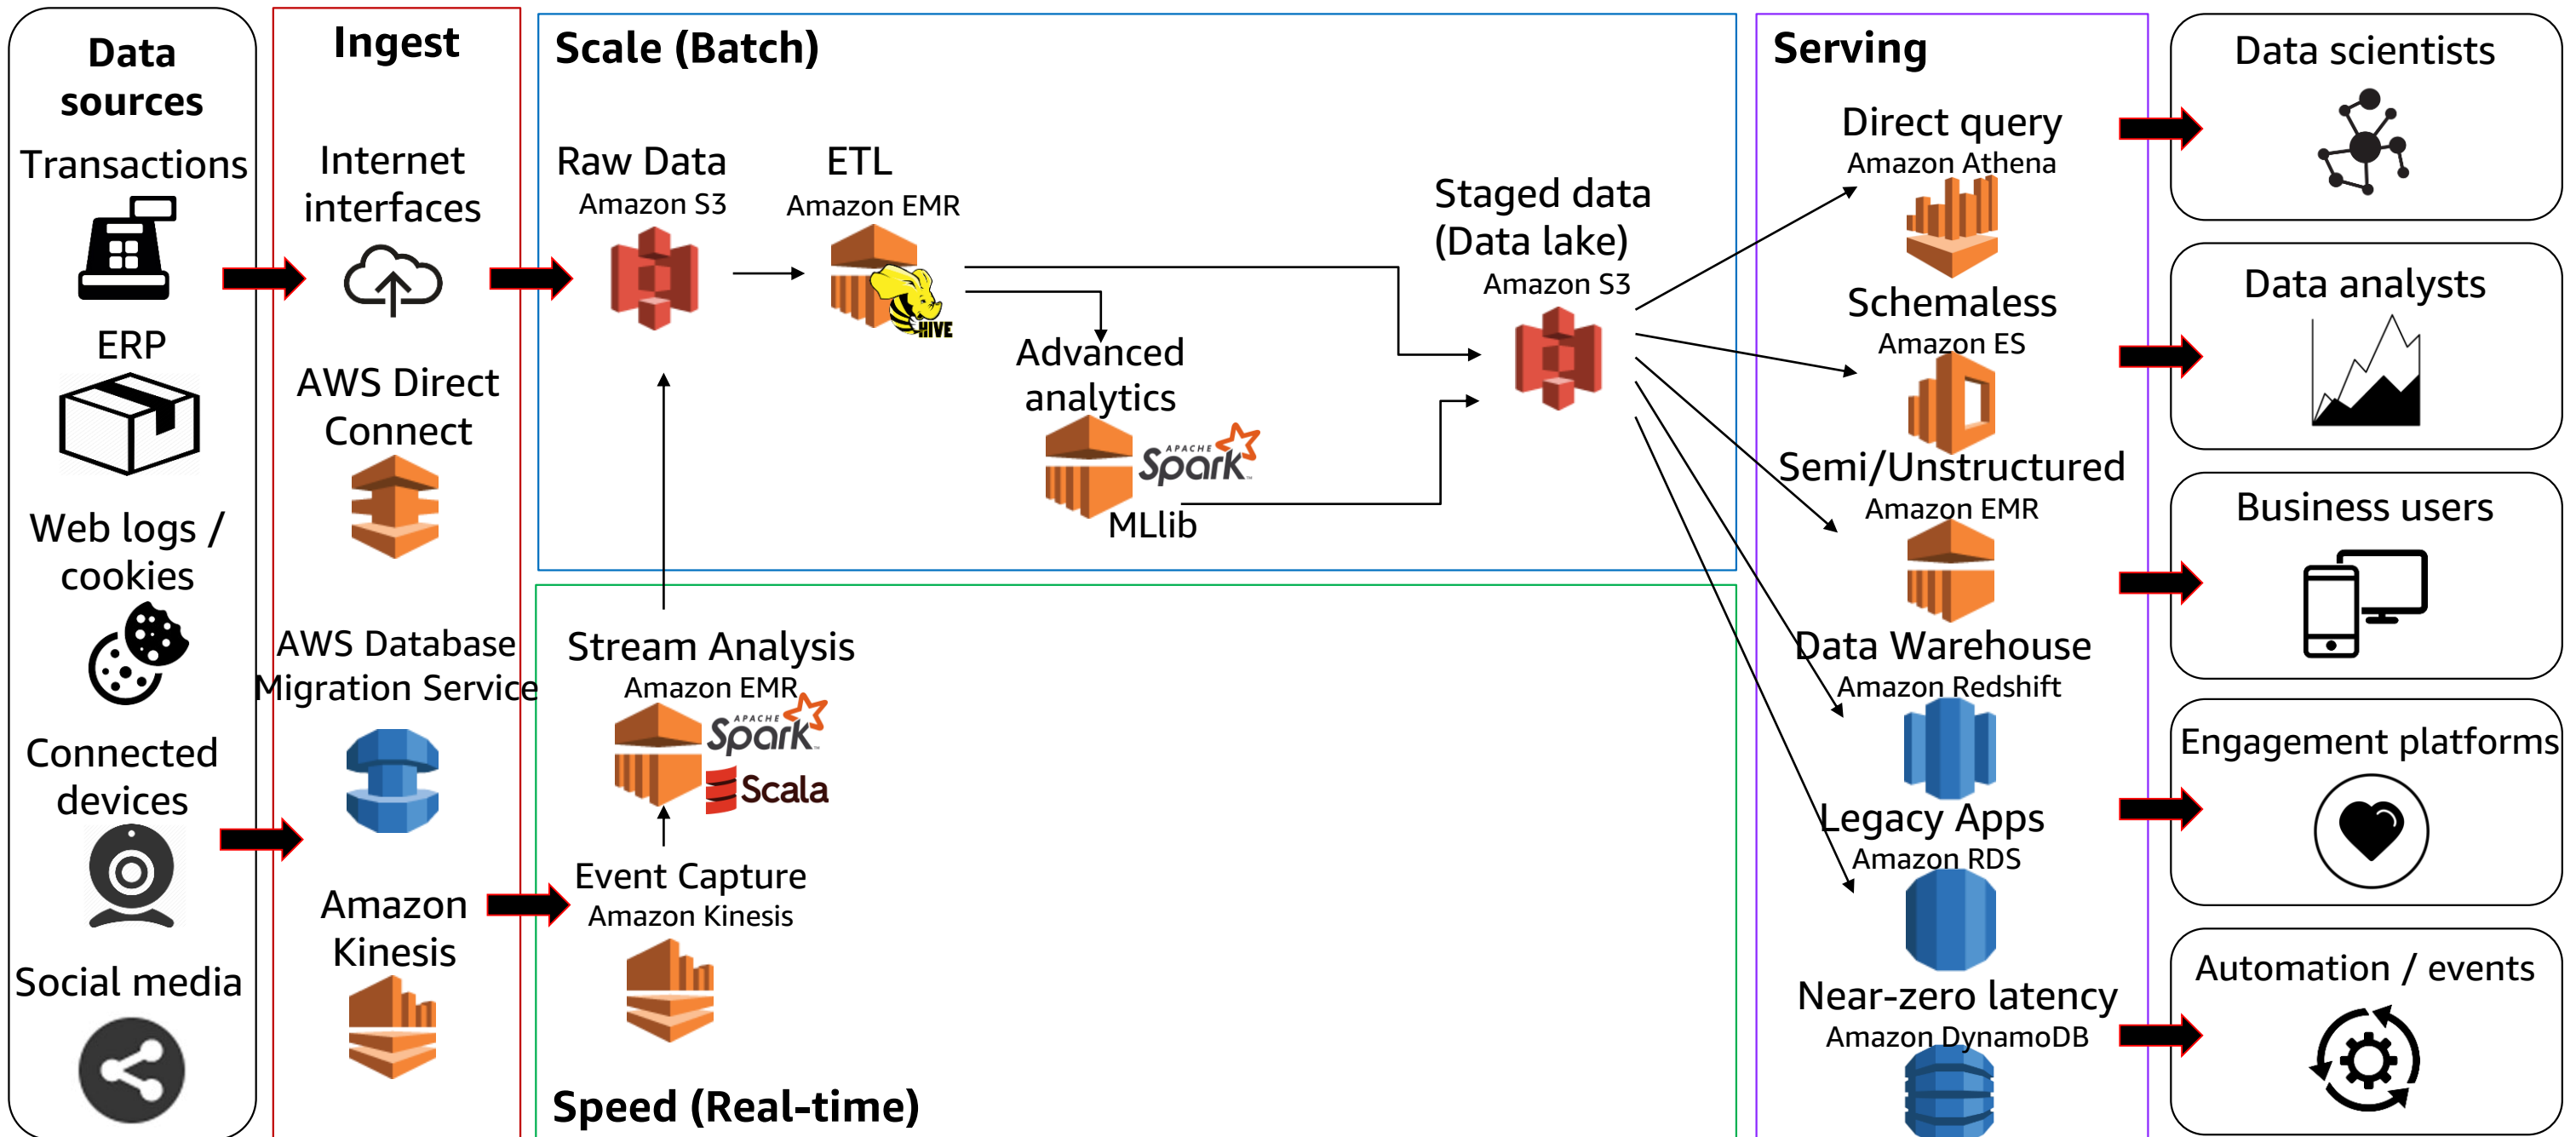
*Kinesis Firehose prepares and loads the data continuously to  
the destinations you chose from among S3, Redshift, Amazon  
Elasticsearch Service, and Kinesis Analytics*

*Analyze streaming data using your favorite BI tools*

- Zero administration and seamless elasticity
- Direct-to-data store integration
- Serverless, continuous data transformations

# Modern data architecture (AWS)

Insights to enhance business applications, new digital services



# Thank you!



Please complete the session  
survey in the mobile app.